

Dog identification

Pongsathorn Utsahawattanasuk 6210554784

Presented to

Assoc.Prof. Somying Thainimit

Dog identification

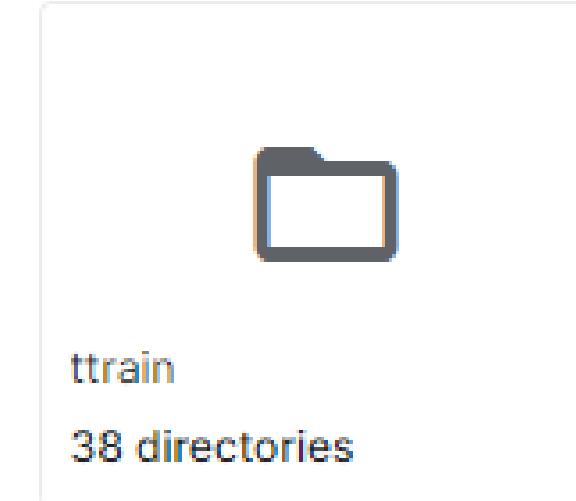
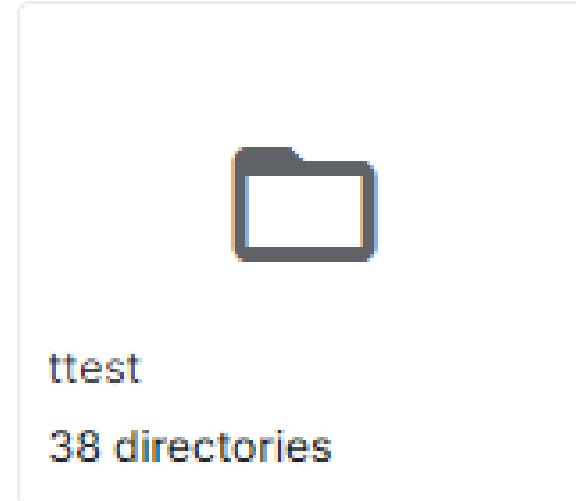
- Dog identification is the process of recognizing and classifying dogs based on their individual features, such as **physical appearance** or behavioral characteristics.
- This process can be useful in a variety of settings, including law enforcement, animal shelters, and scientific research.
- Various methods can be used for dog identification, such as visual identification by trained professionals, microchipping, and DNA testing.
- Advances in technology, such as **machine learning and computer vision**, have made it possible to automate dog identification using image and video data.



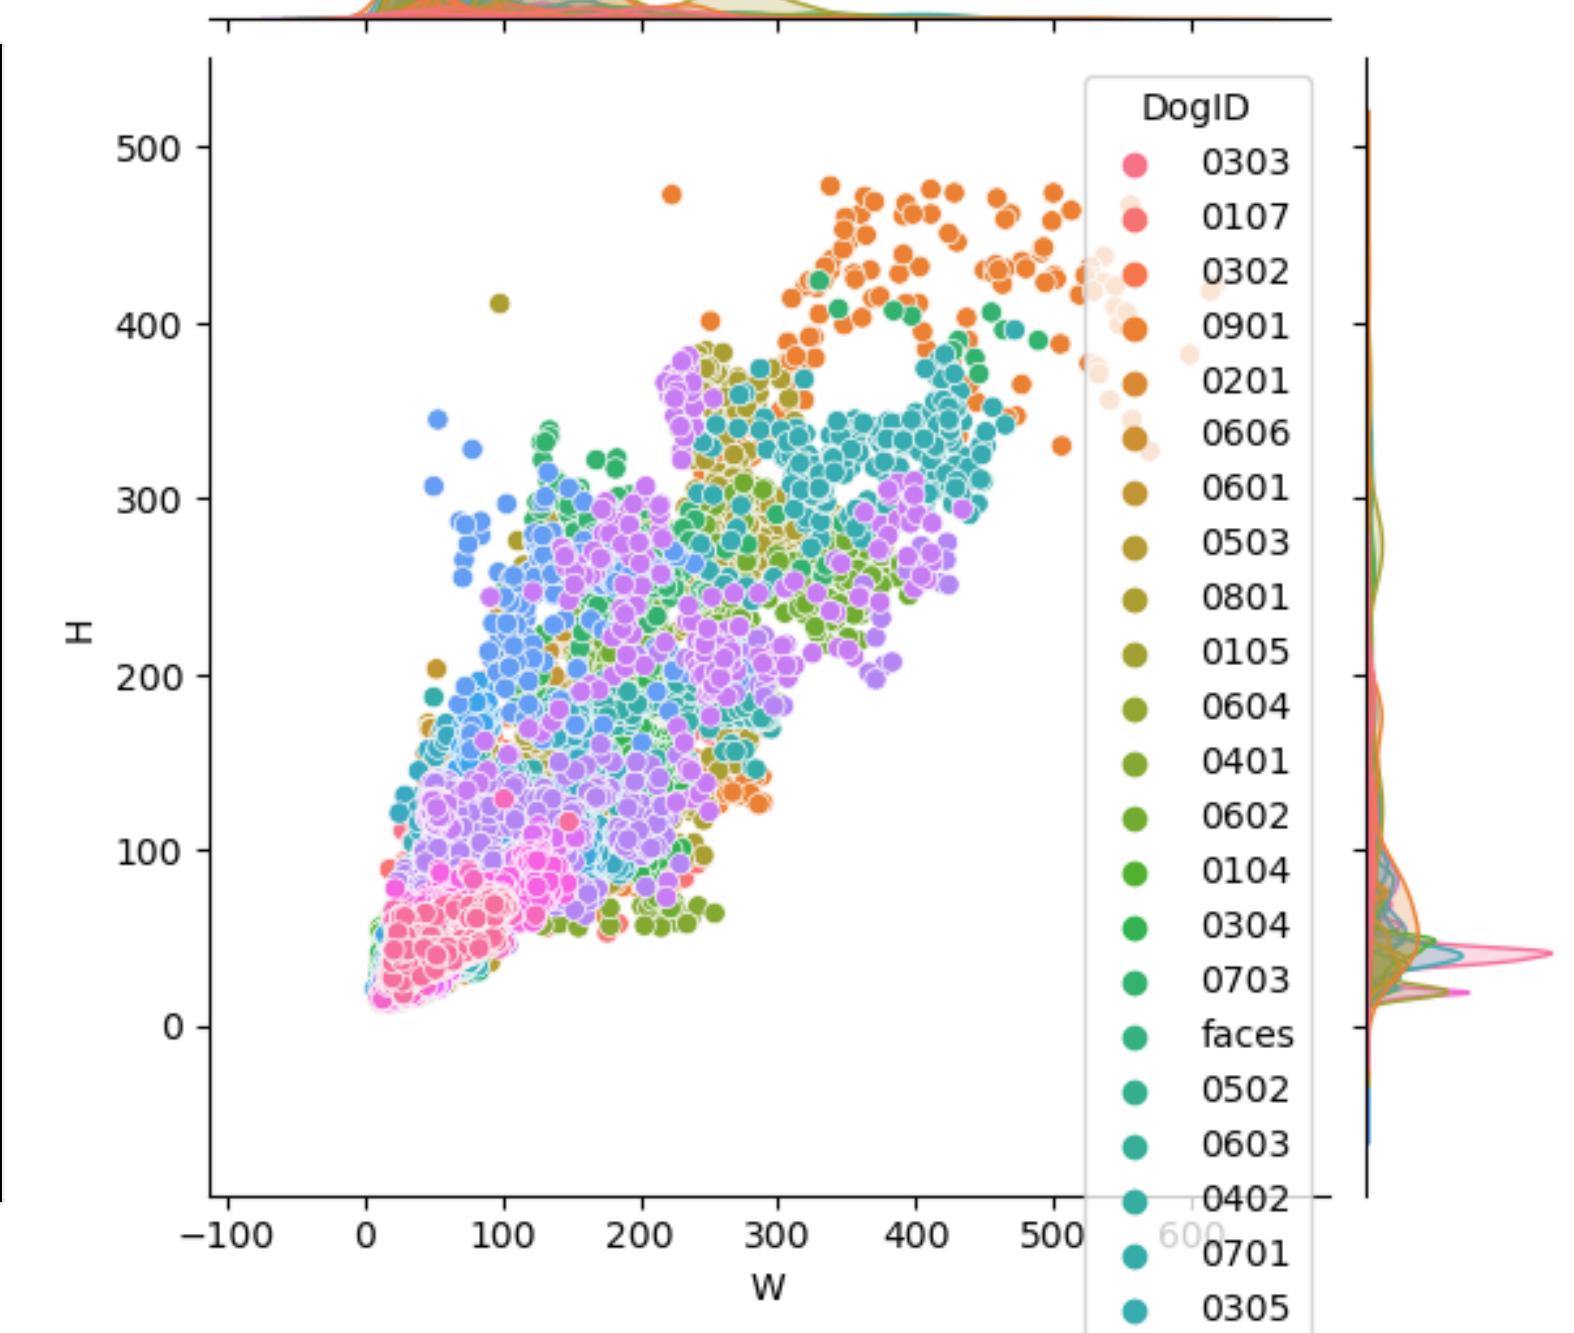
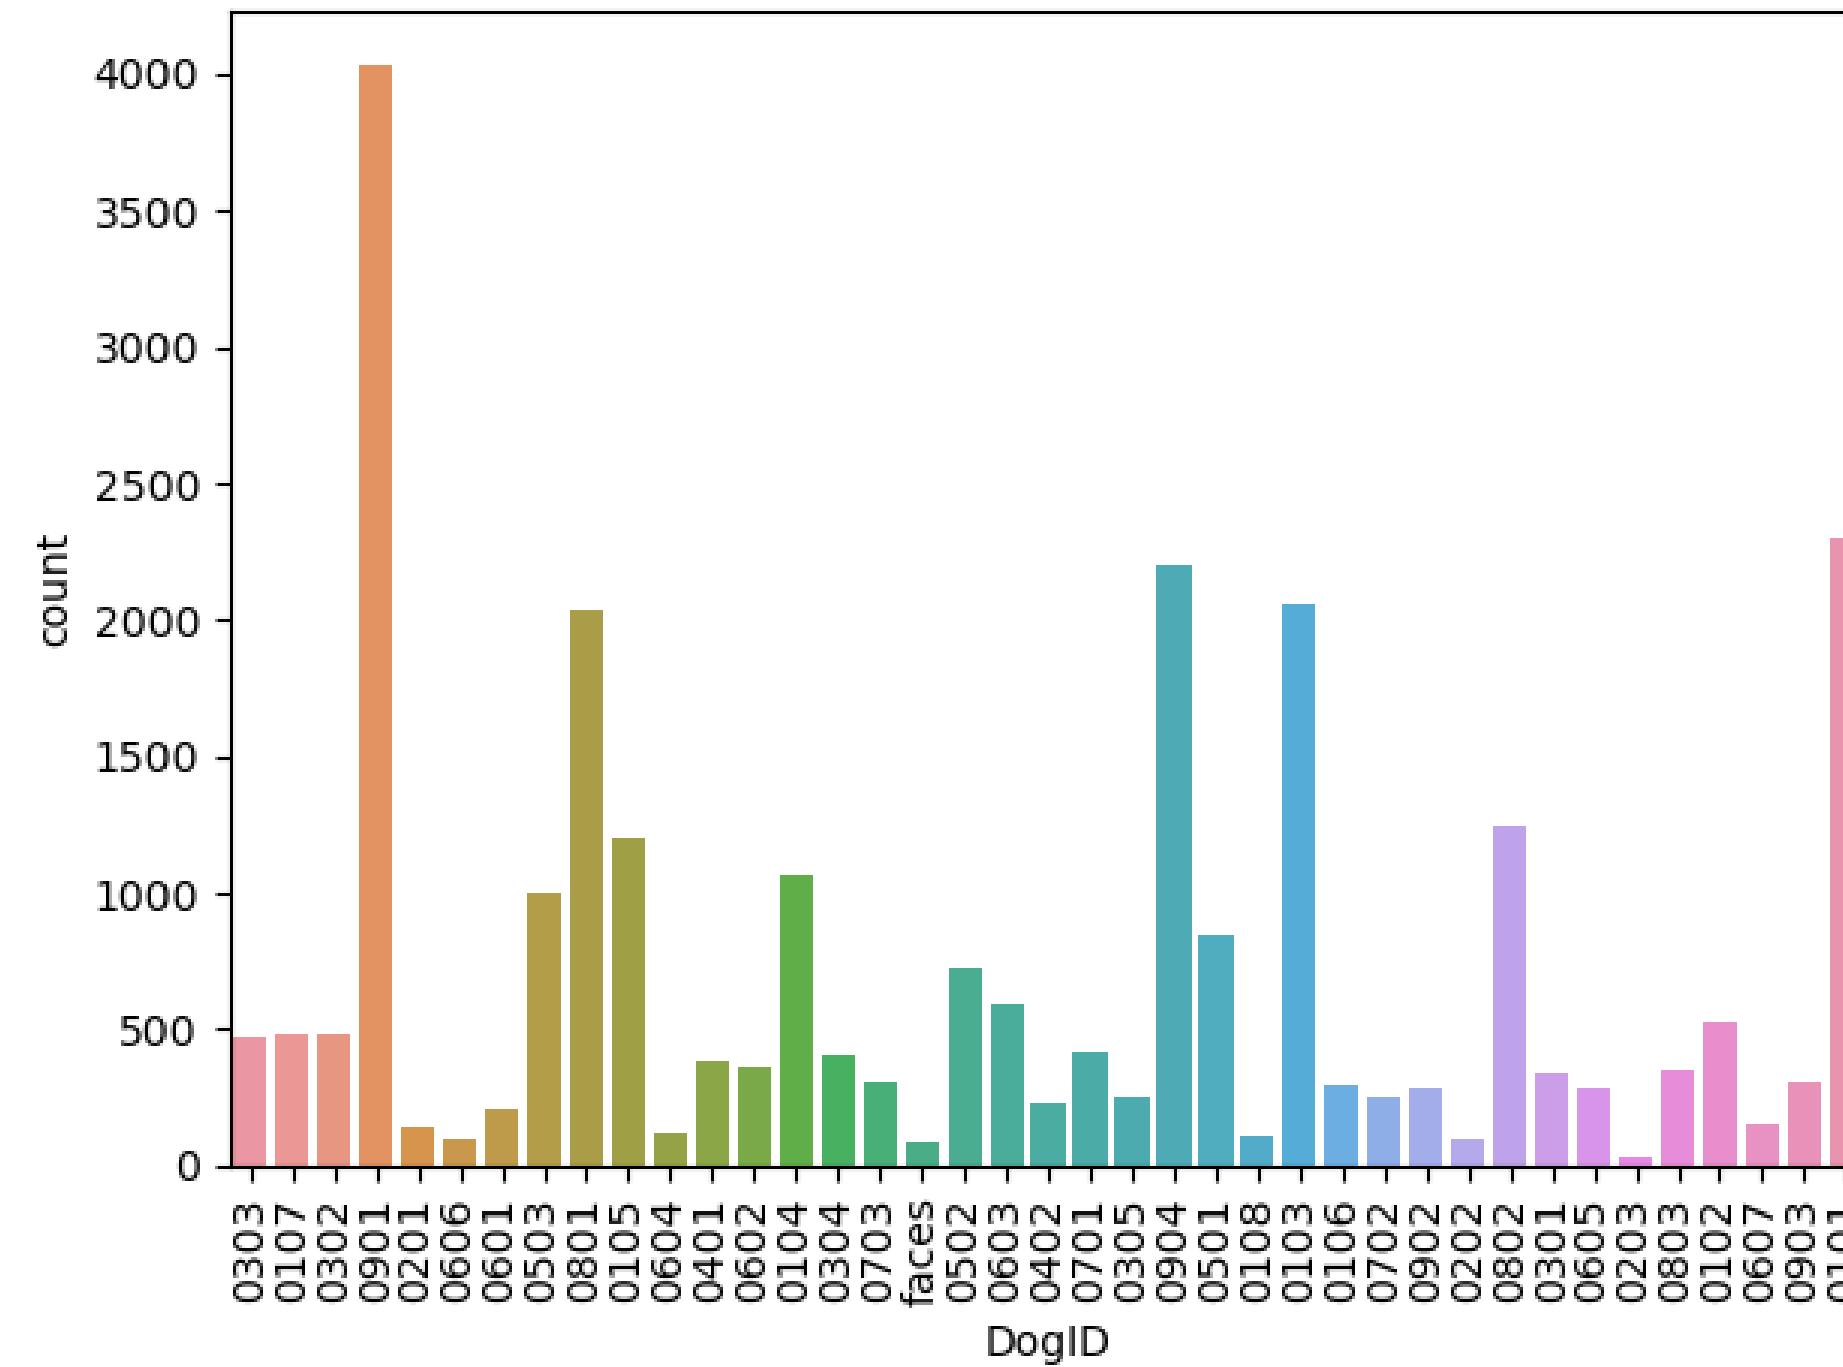
Data

Input (174.57 MB)

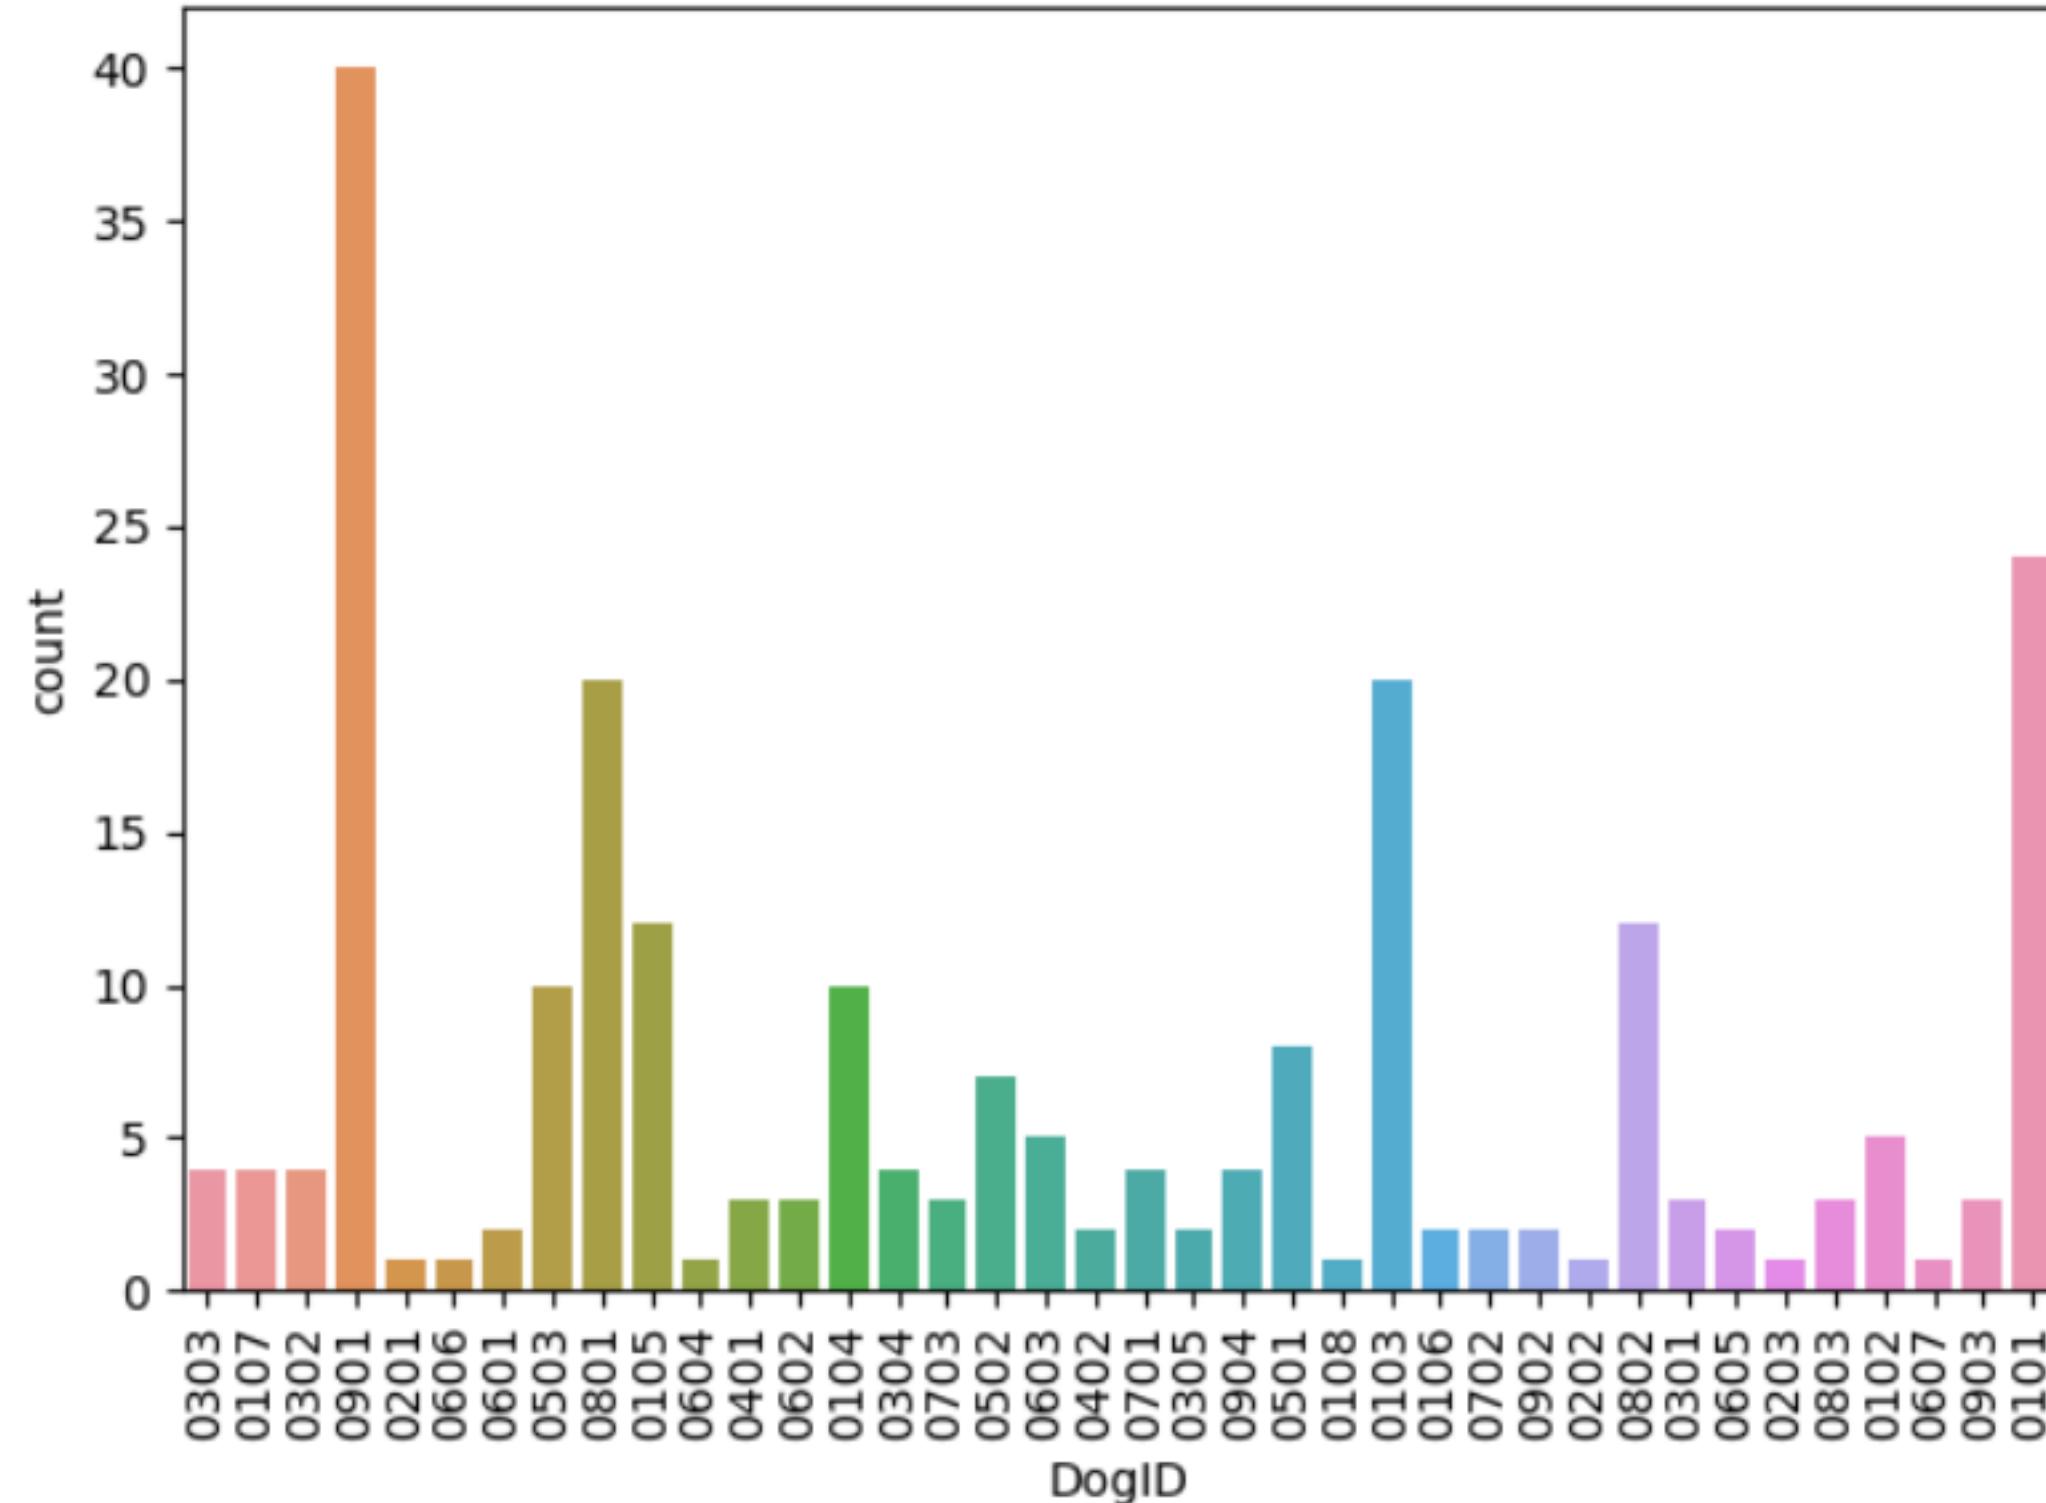
- Data Sources
 -  doggie
 - □ dataset
 - ▶ □ ttest
 - □ ttrain
 - ▶ □ DogID.0101
 - ▶ □ DogID.0102
 - ▶ □ DogID.0103
 - ▶ □ DogID.0104
 - ▶ □ DogID.0105
 - ▶ □ DogID.0106
 - ▶ □ DogID.0107
 - ▶ □ DogID.0108
 - ▶ □ DogID.0201
 - ▶ □ DogID.0202



EDA: train data



EDA: Test data



Methodology

Siamese Network



image 1

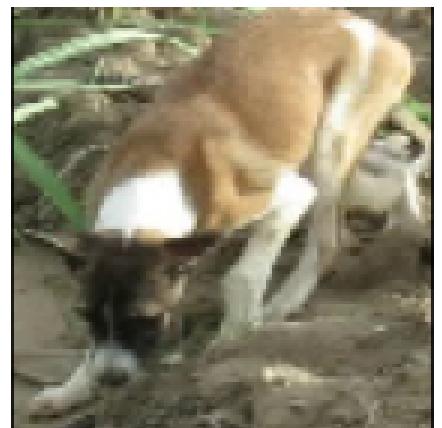
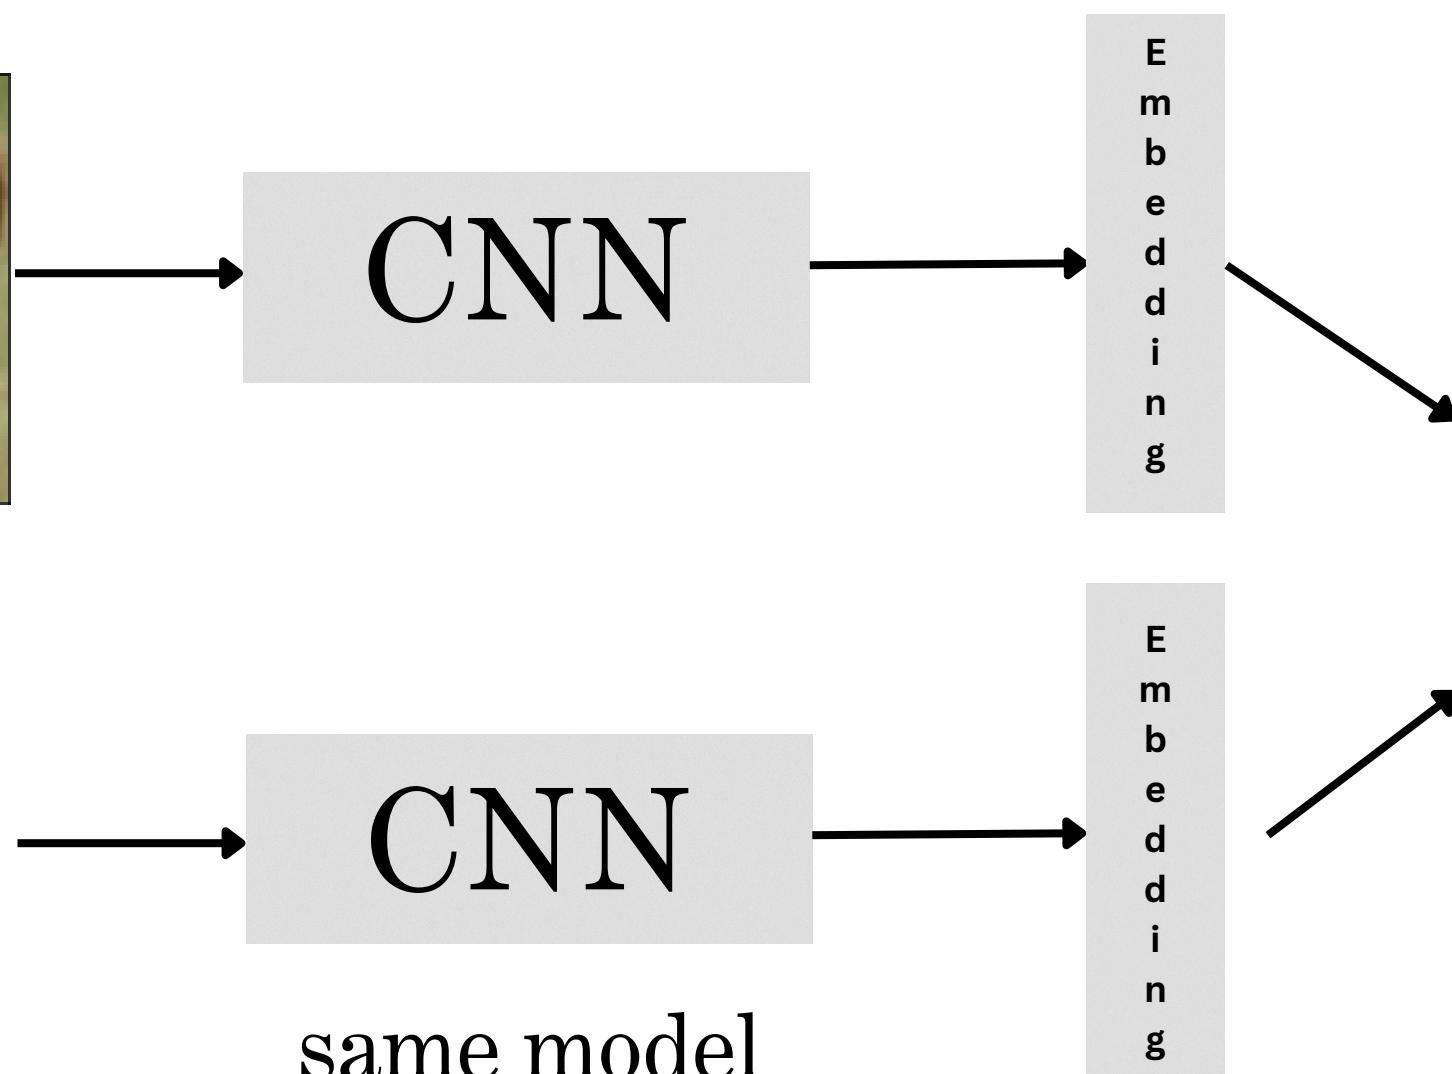


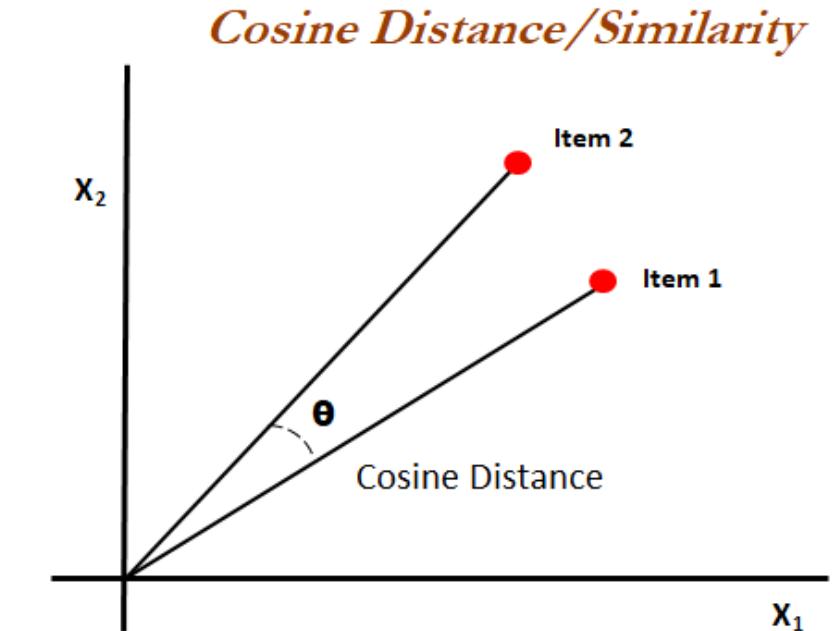
image 2



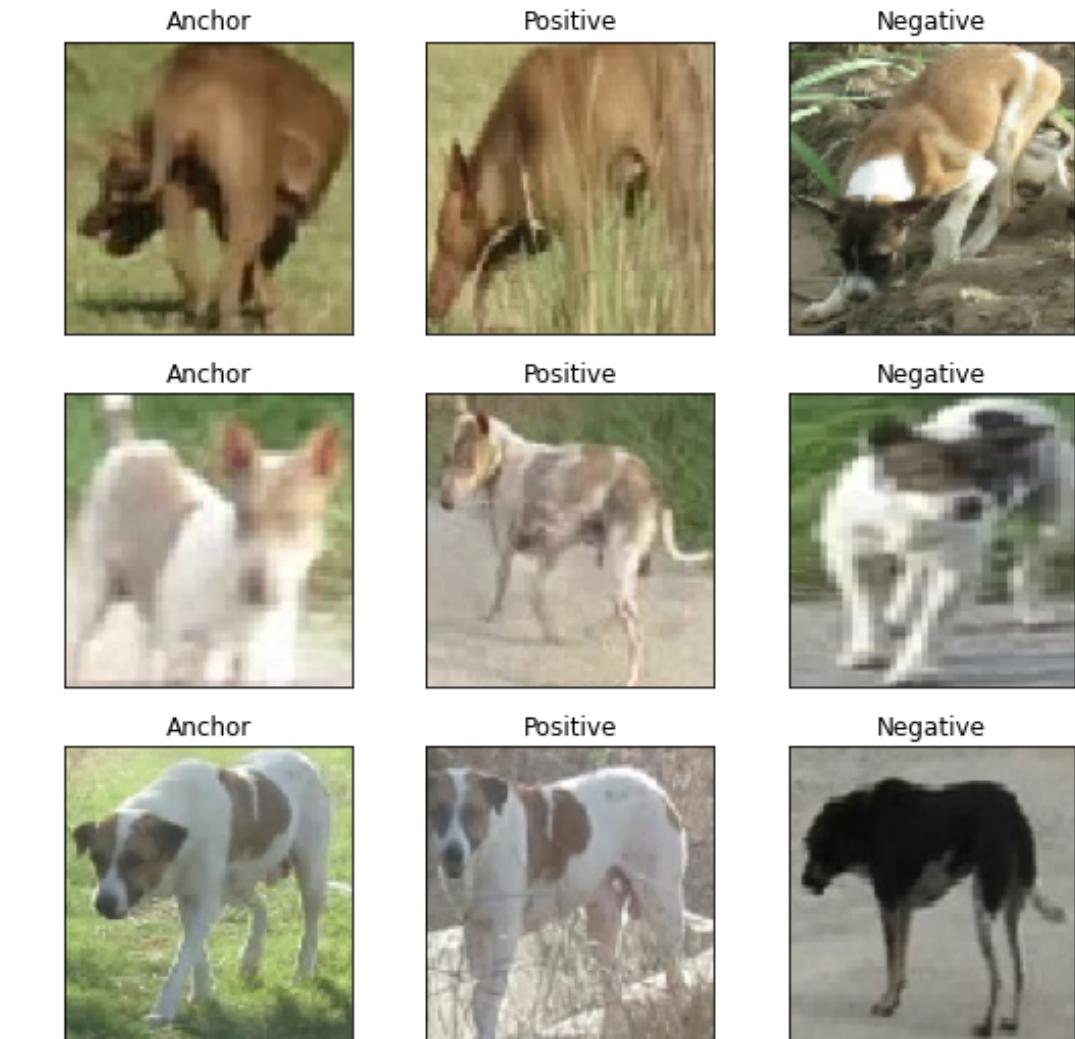
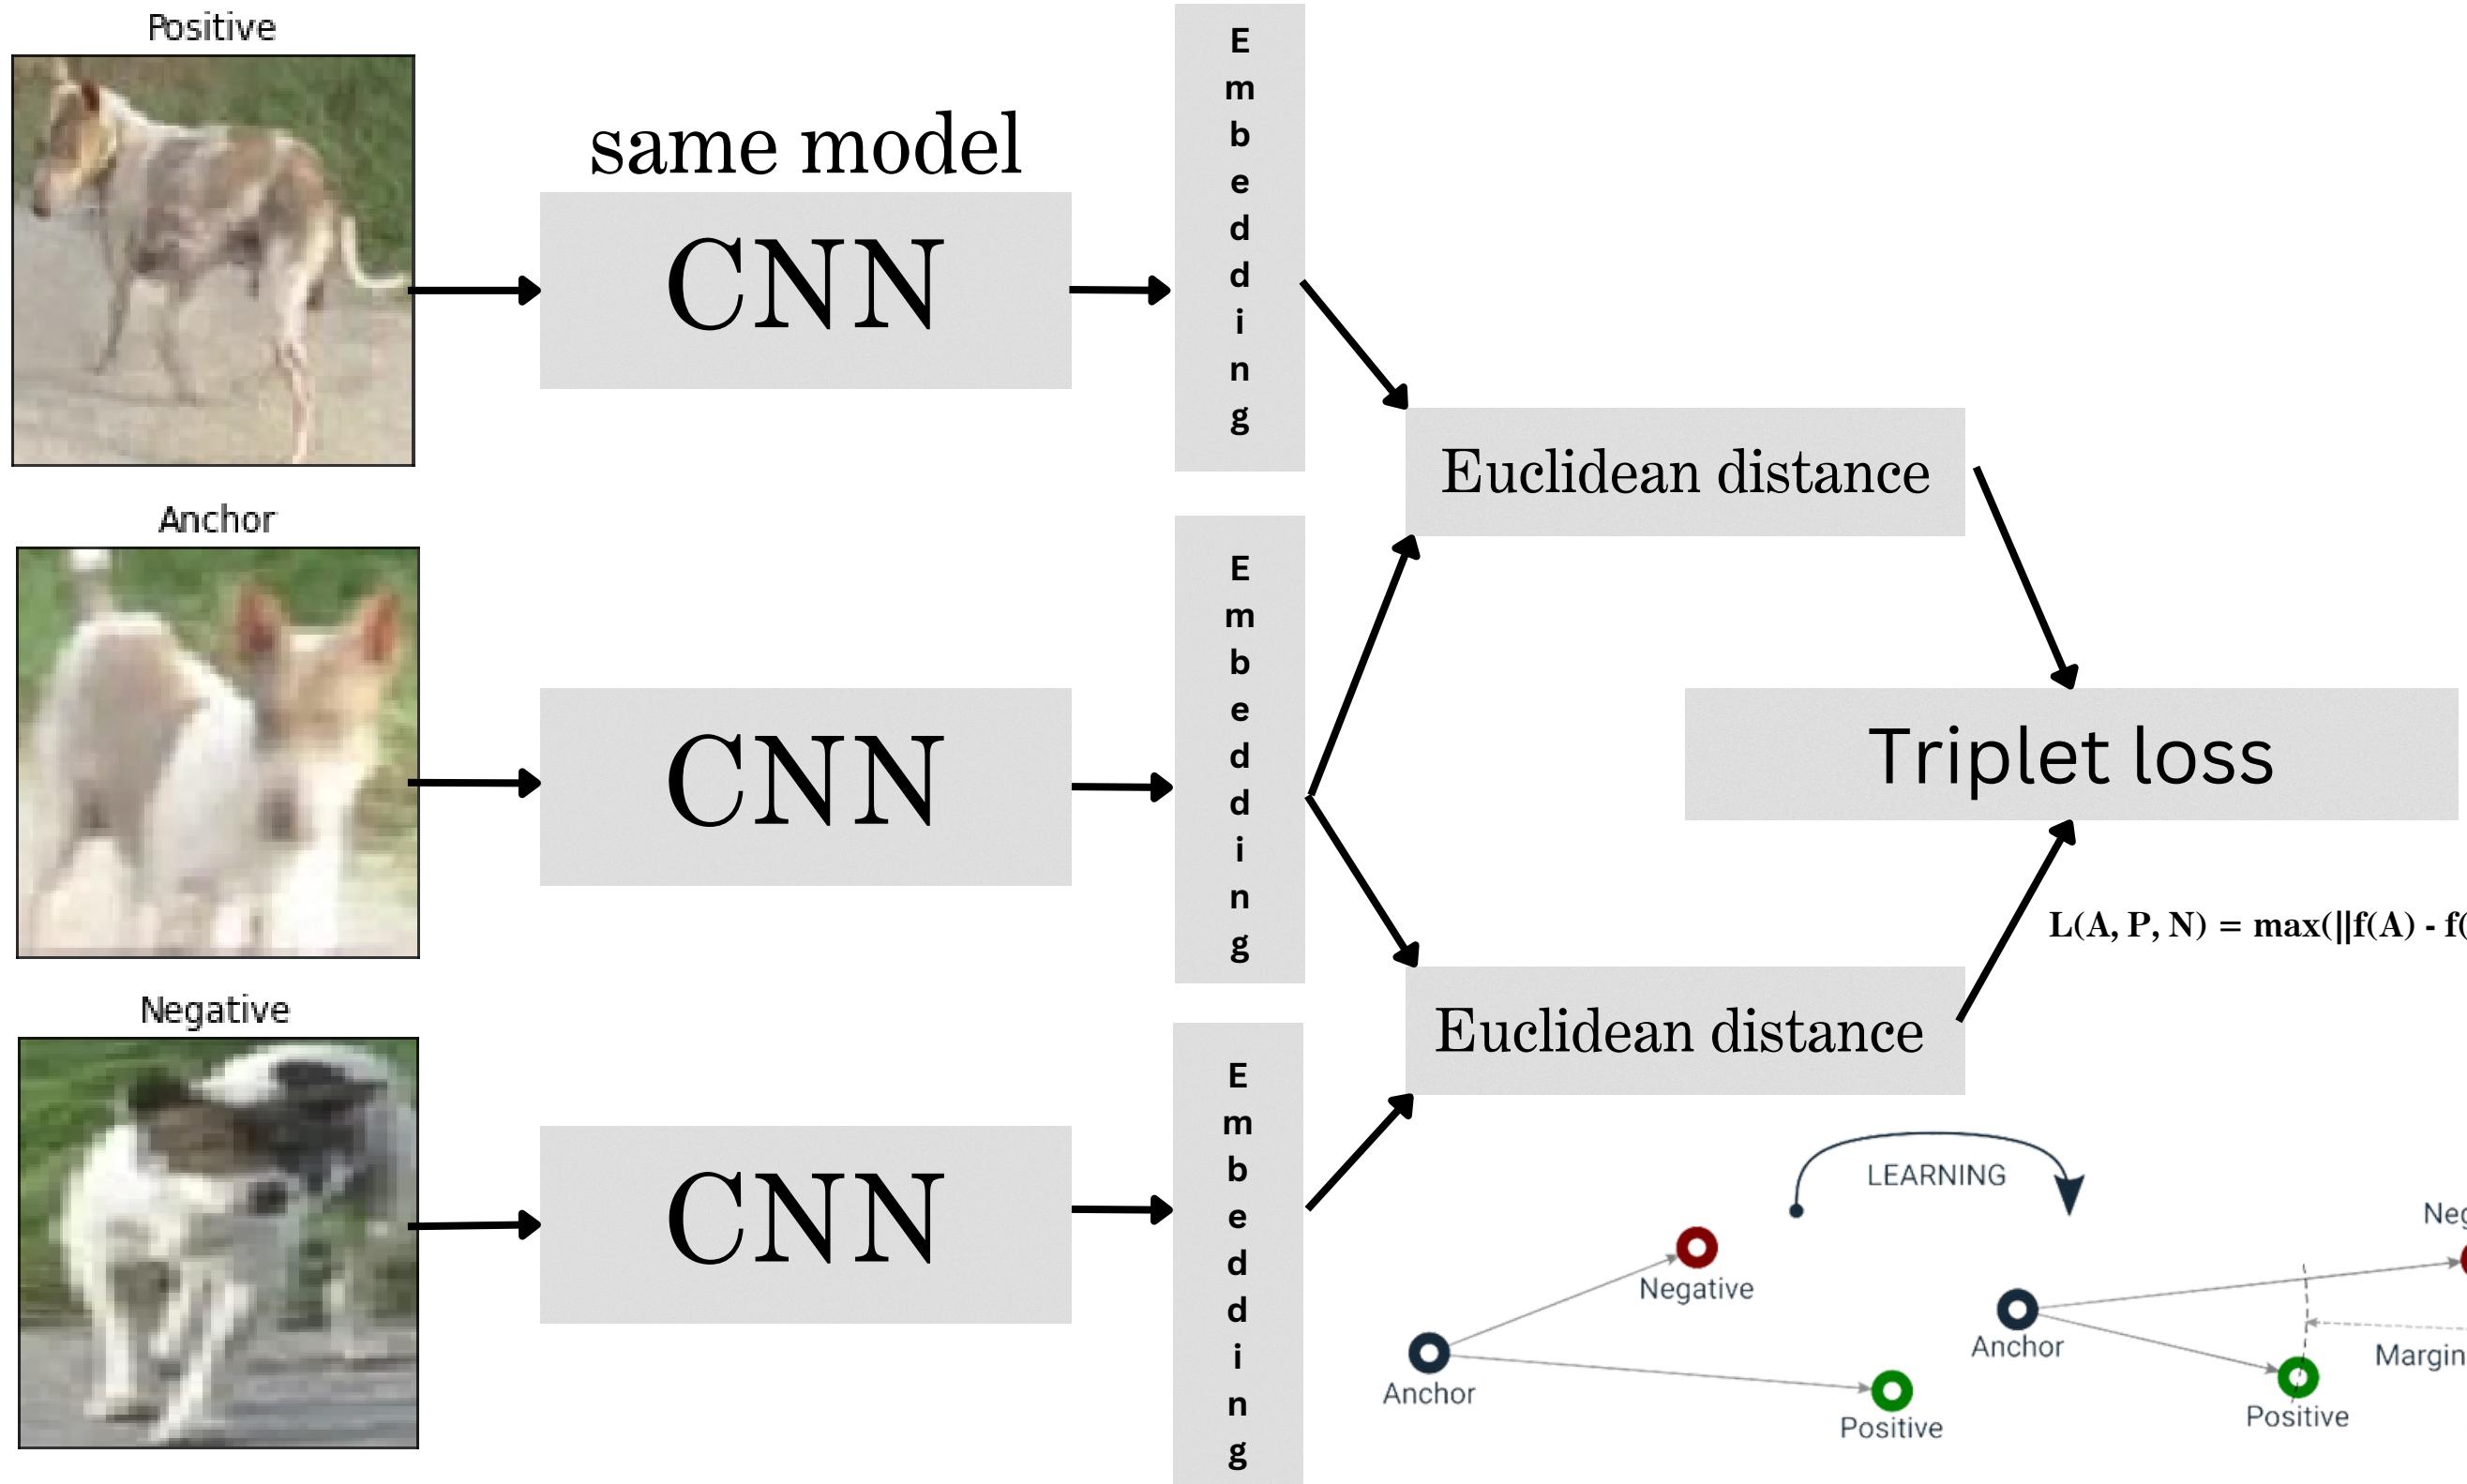
Compute cosine similarity matrix

-1 to 1

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\| \|\mathbf{B}\|} = \frac{\sum_{i=1}^n A_i B_i}{\sqrt{\sum_{i=1}^n A_i^2} \sqrt{\sum_{i=1}^n B_i^2}},$$



My Siamese Network



Keras documentation: Image similarity estimation using a Siamese Network with a triplet loss

Keras documentation

[K keras.io /](https://keras.io/)

Data Generator Code Snippet

	filepath	W	H	DogID
0	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	276	183	0303
1	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	211	193	0303
2	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	193	189	0303
3	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	215	189	0303
4	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	260	182	0303
...
6499	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	60	41	0101
6500	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	51	32	0101
6501	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	31	44	0101
6502	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	43	44	0101
6503	/kaggle/input/dogdataset/Gallery_sameCamera/Do...	83	44	0101

```

def data_generator(data_df, batch_size=BATCH_SIZE):
    while True:
        data_df = data_df.sample(frac=1).reset_index(drop=True) # Shuffle the dataframe
        for i in range(0, len(data_df), batch_size):
            batch_df = data_df[i:i + batch_size]
            anchor_batch, positive_batch, negative_batch = [], [], []
            for _, row in batch_df.iterrows():
                anchor_img_path = row['filepath']
                anchor_dog_id = row['DogID']

                # Get positive image
                positive_df = data_df[data_df['DogID'] == anchor_dog_id].sample(n=2)
                positive_row = positive_df.iloc[0] if positive_df.iloc[0]['filepath'] != anchor_img_path else positive_df.iloc[1]
                positive_img_path = positive_row['filepath']

                # Get negative image
                negative_df = data_df[data_df['DogID'] != anchor_dog_id].sample(1)
                negative_img_path = negative_df.iloc[0]['filepath']

                anchor_batch.append(load_image(anchor_img_path))
                positive_batch.append(load_image(positive_img_path))
                negative_batch.append(load_image(negative_img_path))

            anchor_batch = np.array(anchor_batch)
            positive_batch = np.array(positive_batch)
            negative_batch = np.array(negative_batch)
            dummy_labels = np.zeros((batch_size, 1)) # Add dummy labels
            yield [anchor_batch, positive_batch, negative_batch], dummy_labels
            gc.collect()
  
```

Pretrained Models

```
# Constants
IMG_SIZE = int(df.describe()['W']['mean']),int(df.describe()['H']['mean'])
BATCH_SIZE = 16
NUM_CLASSES = len(df['DogID'].unique())
print(IMG_SIZE)
```

☞ (93, 91)

Choose Pretrained Model



model_name:

VGG19

VGG16

VGG19

ResNet50

InceptionV3

InceptionResNetV2

MobileNet

MobileNetV2

DenseNet121

DenseNet169

EfficientNetB0

```
[ ] # Add dense layer
x = layers.GlobalAveragePooling2D()(x)
x = layers.Dense(128)(x)
output = layers.Dense(128)(x)
```

```
# Create an embedding model with the input and output of the base CNN model
embedding_model = Model(inputs=base_cnn.input, outputs=output)
```

My Siamese Network (VGG19)

```
def create_siamese_model(embedding_model):
    anchor_input = layers.Input(shape=(*IMG_SIZE, 3))
    positive_input = layers.Input(shape=(*IMG_SIZE, 3))
    negative_input = layers.Input(shape=(*IMG_SIZE, 3))

    anchor_embedding = embedding_model(anchor_input)
    positive_embedding = embedding_model(positive_input)
    negative_embedding = embedding_model(negative_input)

    positive_distance = layers.Lambda(euclidean_distance)([anchor_embedding, positive_embedding])
    negative_distance = layers.Lambda(euclidean_distance)([anchor_embedding, negative_embedding])

    stacked_dists = layers.concatenate([positive_distance, negative_distance])
    siamese_model = Model([anchor_input, positive_input, negative_input], stacked_dists)

    siamese_model.compile(optimizer=tf.keras.optimizers.Adam(1e-5), loss=triplet_loss, metrics=[triplet_accuracy])
    return siamese_model
```

My Siamese Network (VGG19)

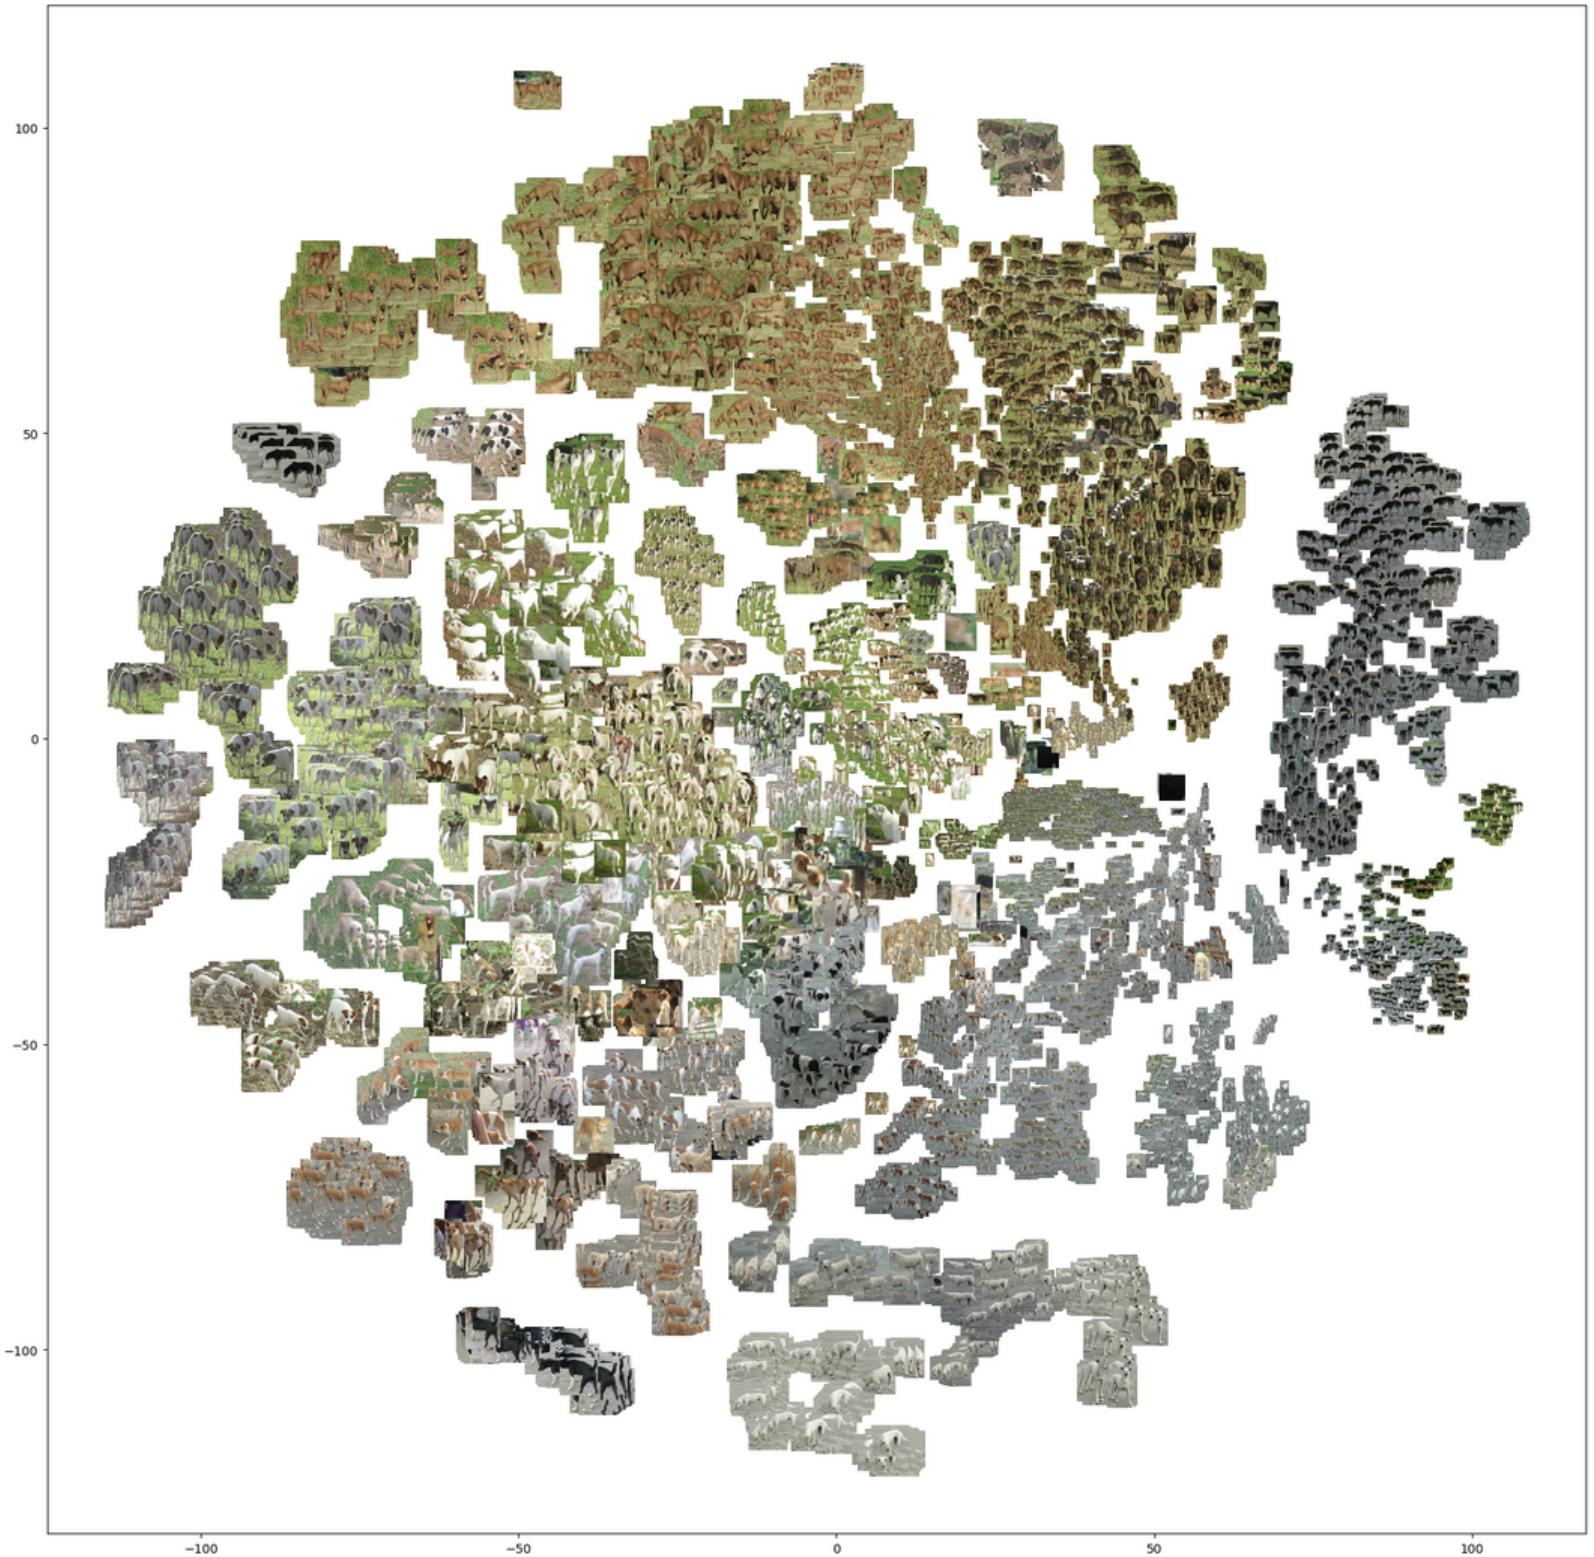
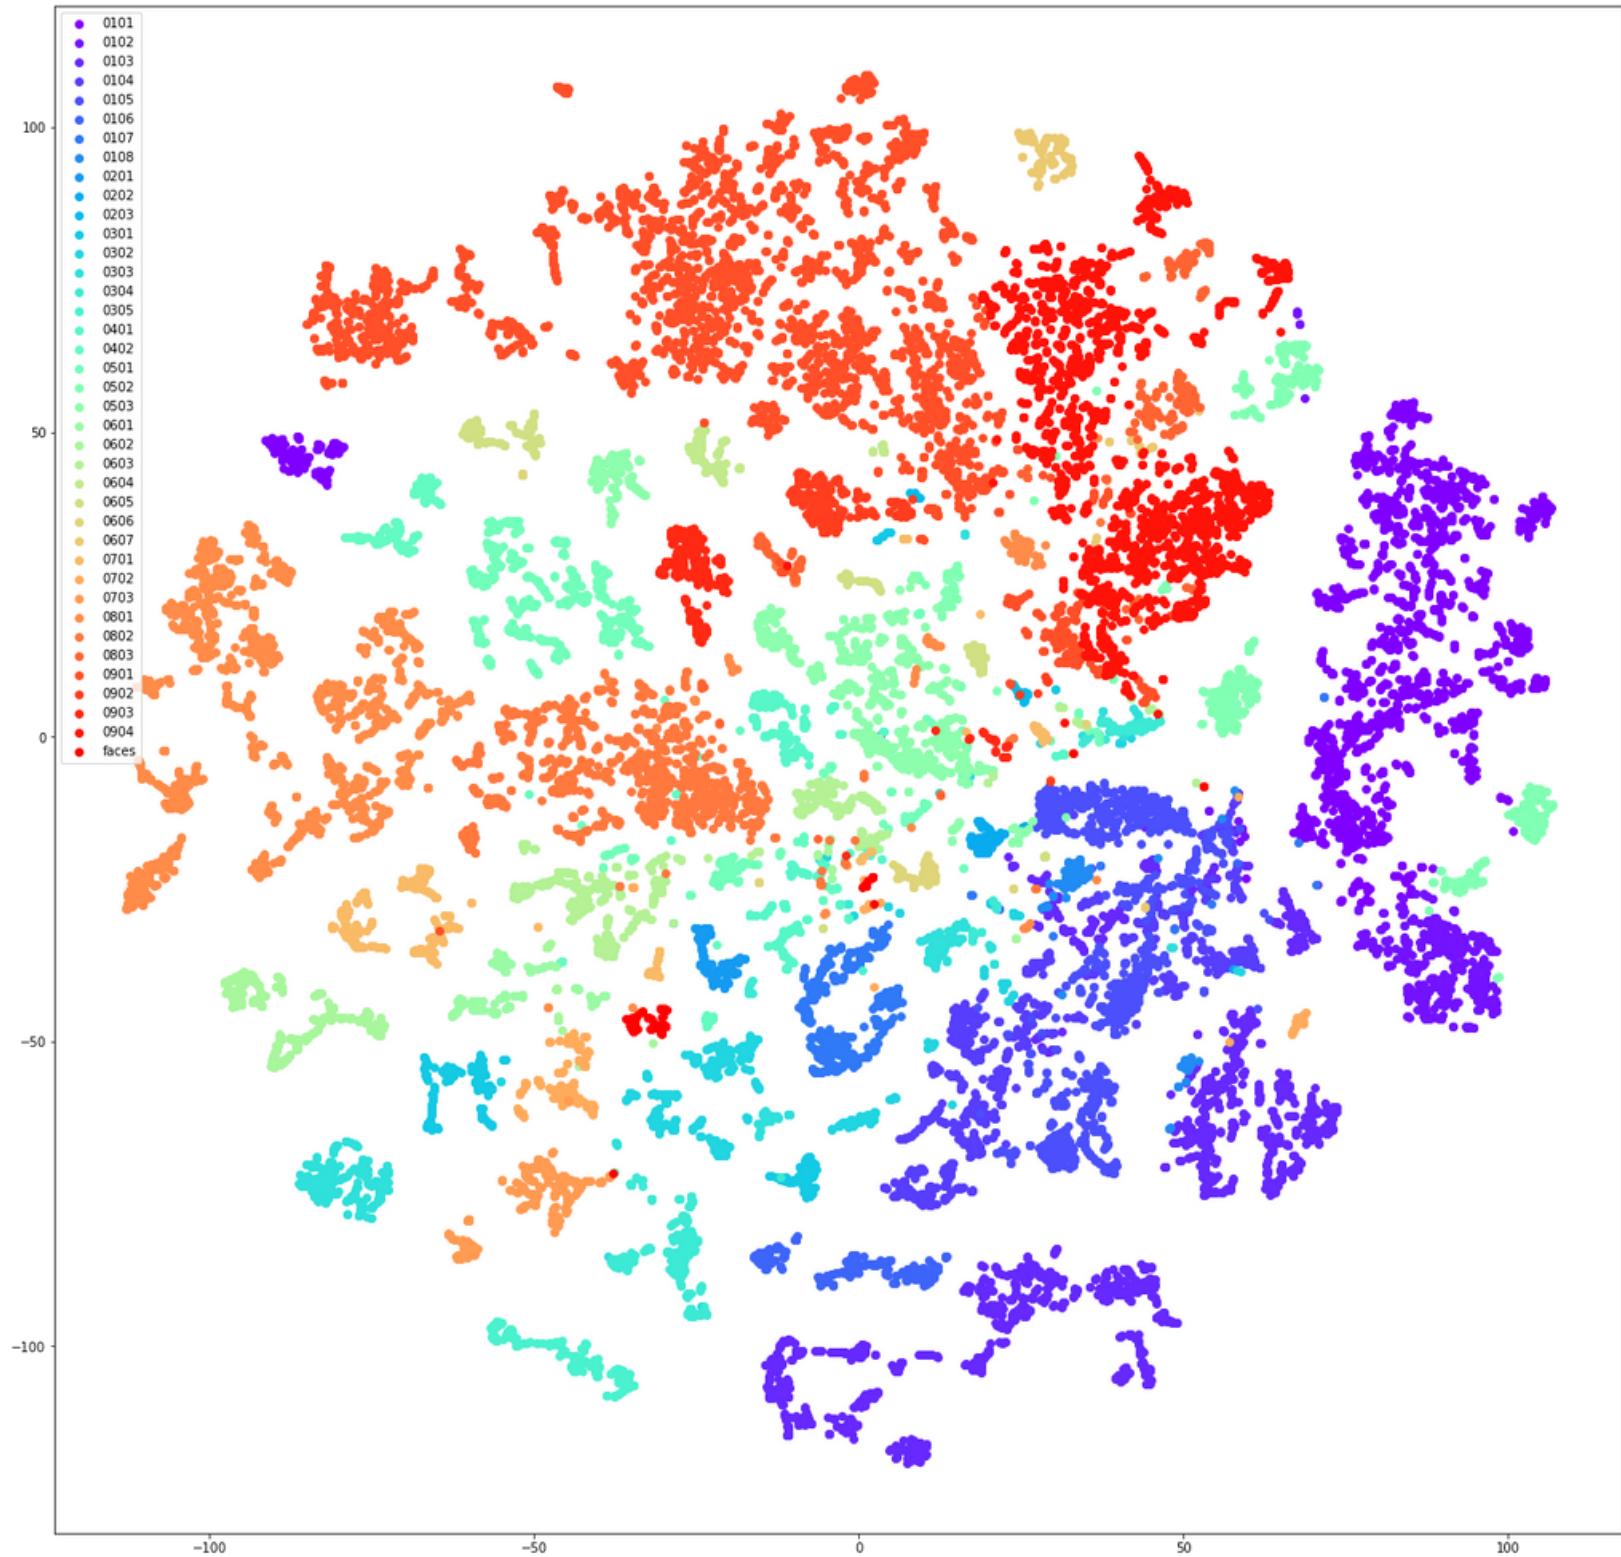
Model: "model_1"

Layer (type)	Output Shape	Param #	Connected to
<hr/>			
input_2 (InputLayer)	[None, 93, 91, 3]	0	[]
input_3 (InputLayer)	[None, 93, 91, 3]	0	[]
input_4 (InputLayer)	[None, 93, 91, 3]	0	[]
model (Functional)	(None, 128)	20454336	['input_2[0][0]', 'input_3[0][0]', 'input_4[0][0]']
lambda (Lambda)	(None, 1)	0	['model[0][0]', 'model[1][0]']
lambda_1 (Lambda)	(None, 1)	0	['model[0][0]', 'model[2][0]']
concatenate (Concatenate)	(None, 2)	0	['lambda[0][0]', 'lambda_1[0][0]']
<hr/>			
Total params:	20,454,336		
Trainable params:	428,416		
Non-trainable params:	20,025,920		

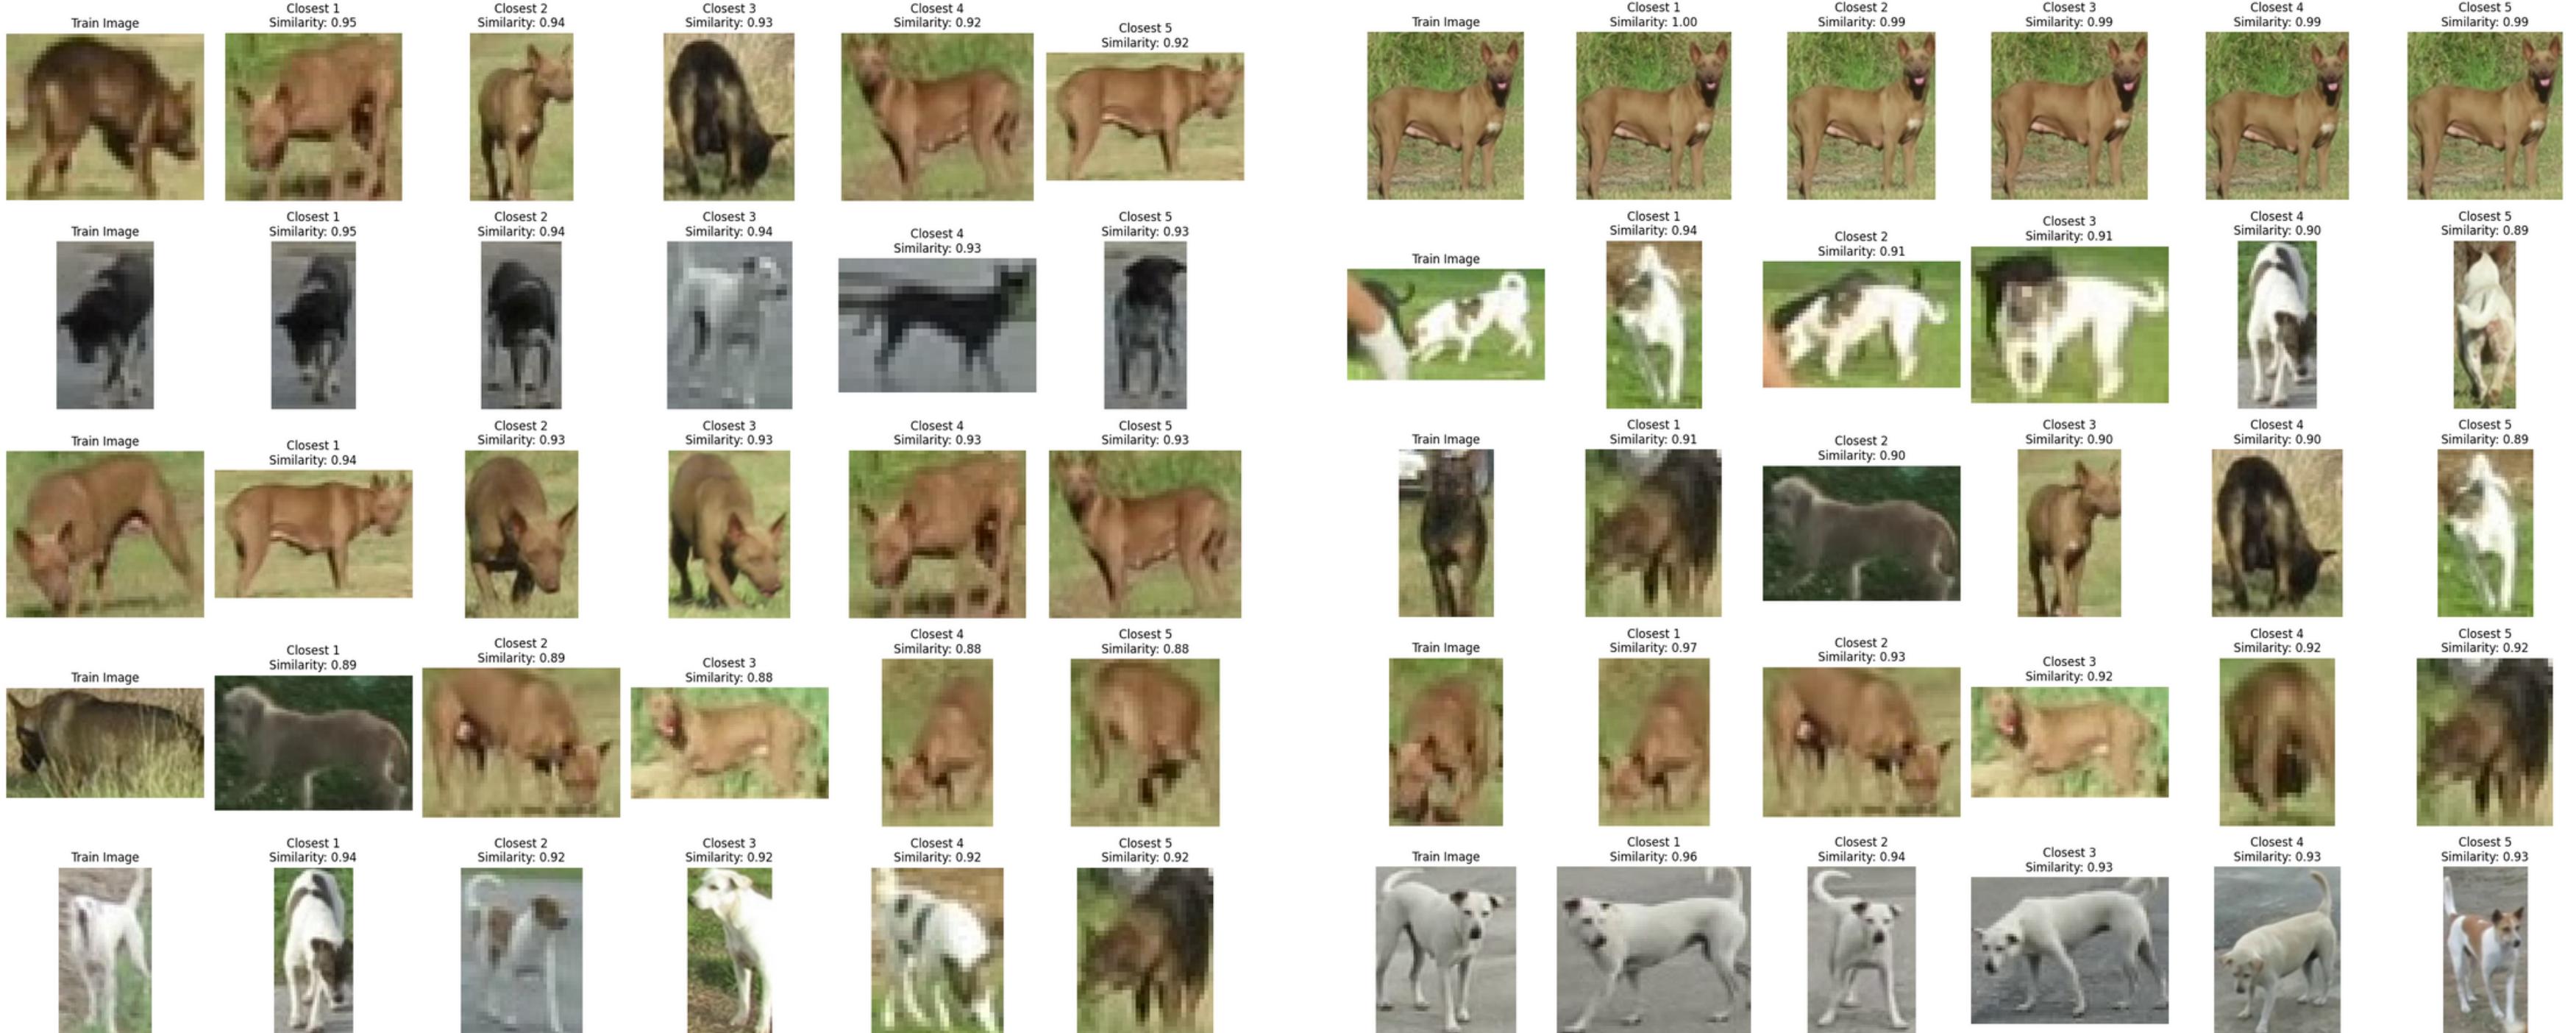
output is concat of positve and
negative distance

Epoch 1/20
 1341/1341 [=====] - 291s 205ms/step - loss: 0.3177 - val_loss: 0.2324
 Epoch 2/20
 1341/1341 [=====] - 275s 205ms/step - loss: 0.1453 - val_loss: 0.1696
 Epoch 3/20
 1341/1341 [=====] - 271s 202ms/step - loss: 0.1125 - val_loss: 0.1330
 Epoch 4/20
 1341/1341 [=====] - 264s 197ms/step - loss: 0.0980 - val_loss: 0.1391
 Epoch 5/20
 1341/1341 [=====] - 280s 209ms/step - loss: 0.0895 - val_loss: 0.0942
 Epoch 6/20
 1341/1341 [=====] - 265s 198ms/step - loss: 0.0732 - val_loss: 0.0876
 Epoch 7/20
 1341/1341 [=====] - 273s 204ms/step - loss: 0.0723 - val_loss: 0.1010
 Epoch 8/20
 1341/1341 [=====] - 273s 204ms/step - loss: 0.0680 - val_loss: 0.0873
 Epoch 9/20
 1341/1341 [=====] - 264s 197ms/step - loss: 0.0568 - val_loss: 0.0803
 Epoch 10/20
 1341/1341 [=====] - 270s 202ms/step - loss: 0.0540 - val_loss: 0.0763
 Epoch 11/20
 1341/1341 [=====] - 267s 199ms/step - loss: 0.0583 - val_loss: 0.0953
 Epoch 12/20
 1341/1341 [=====] - 270s 202ms/step - loss: 0.0535 - val_loss: 0.0730
 Epoch 13/20
 1341/1341 [=====] - 268s 200ms/step - loss: 0.0534 - val_loss: 0.0706
 Epoch 14/20
 1341/1341 [=====] - 268s 200ms/step - loss: 0.0437 - val_loss: 0.0657
 Epoch 15/20
 1341/1341 [=====] - 257s 192ms/step - loss: 0.0456 - val_loss: 0.0696
 Epoch 16/20
 1341/1341 [=====] - 270s 201ms/step - loss: 0.0428 - val_loss: 0.0604
 Epoch 17/20
 1341/1341 [=====] - 272s 203ms/step - loss: 0.0434 - val_loss: 0.0515
 Epoch 18/20
 1341/1341 [=====] - 268s 200ms/step - loss: 0.0382 - val_loss: 0.0616
 Epoch 19/20
 1341/1341 [=====] - 280s 209ms/step - loss: 0.0391 - val_loss: 0.0572
 Epoch 20/20
 1341/1341 [=====] - 272s 203ms/step - loss: 0.0375 - val_loss: 0.0533

Train's Embedding features

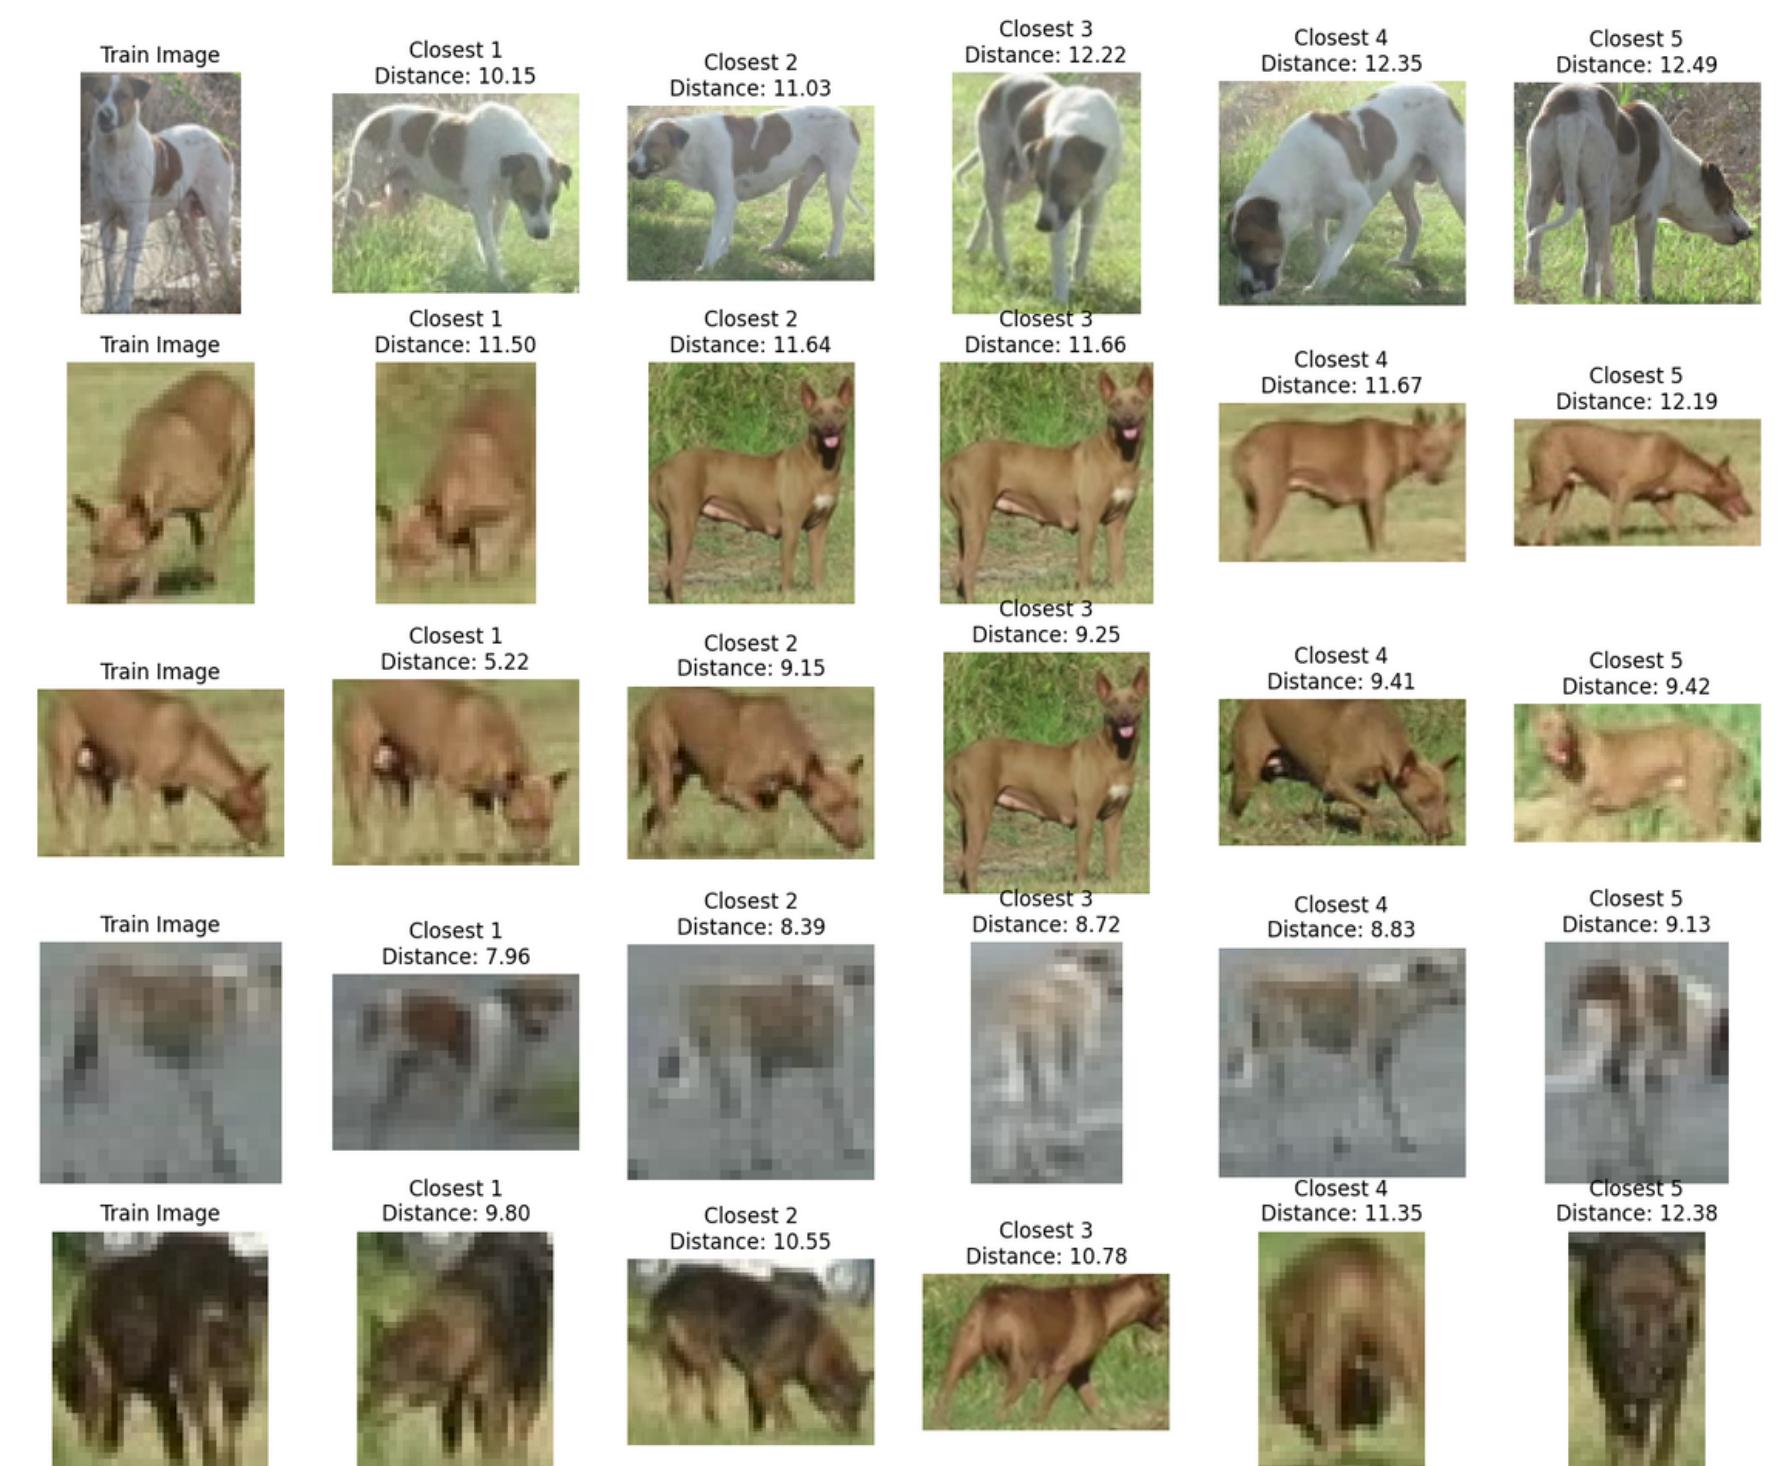
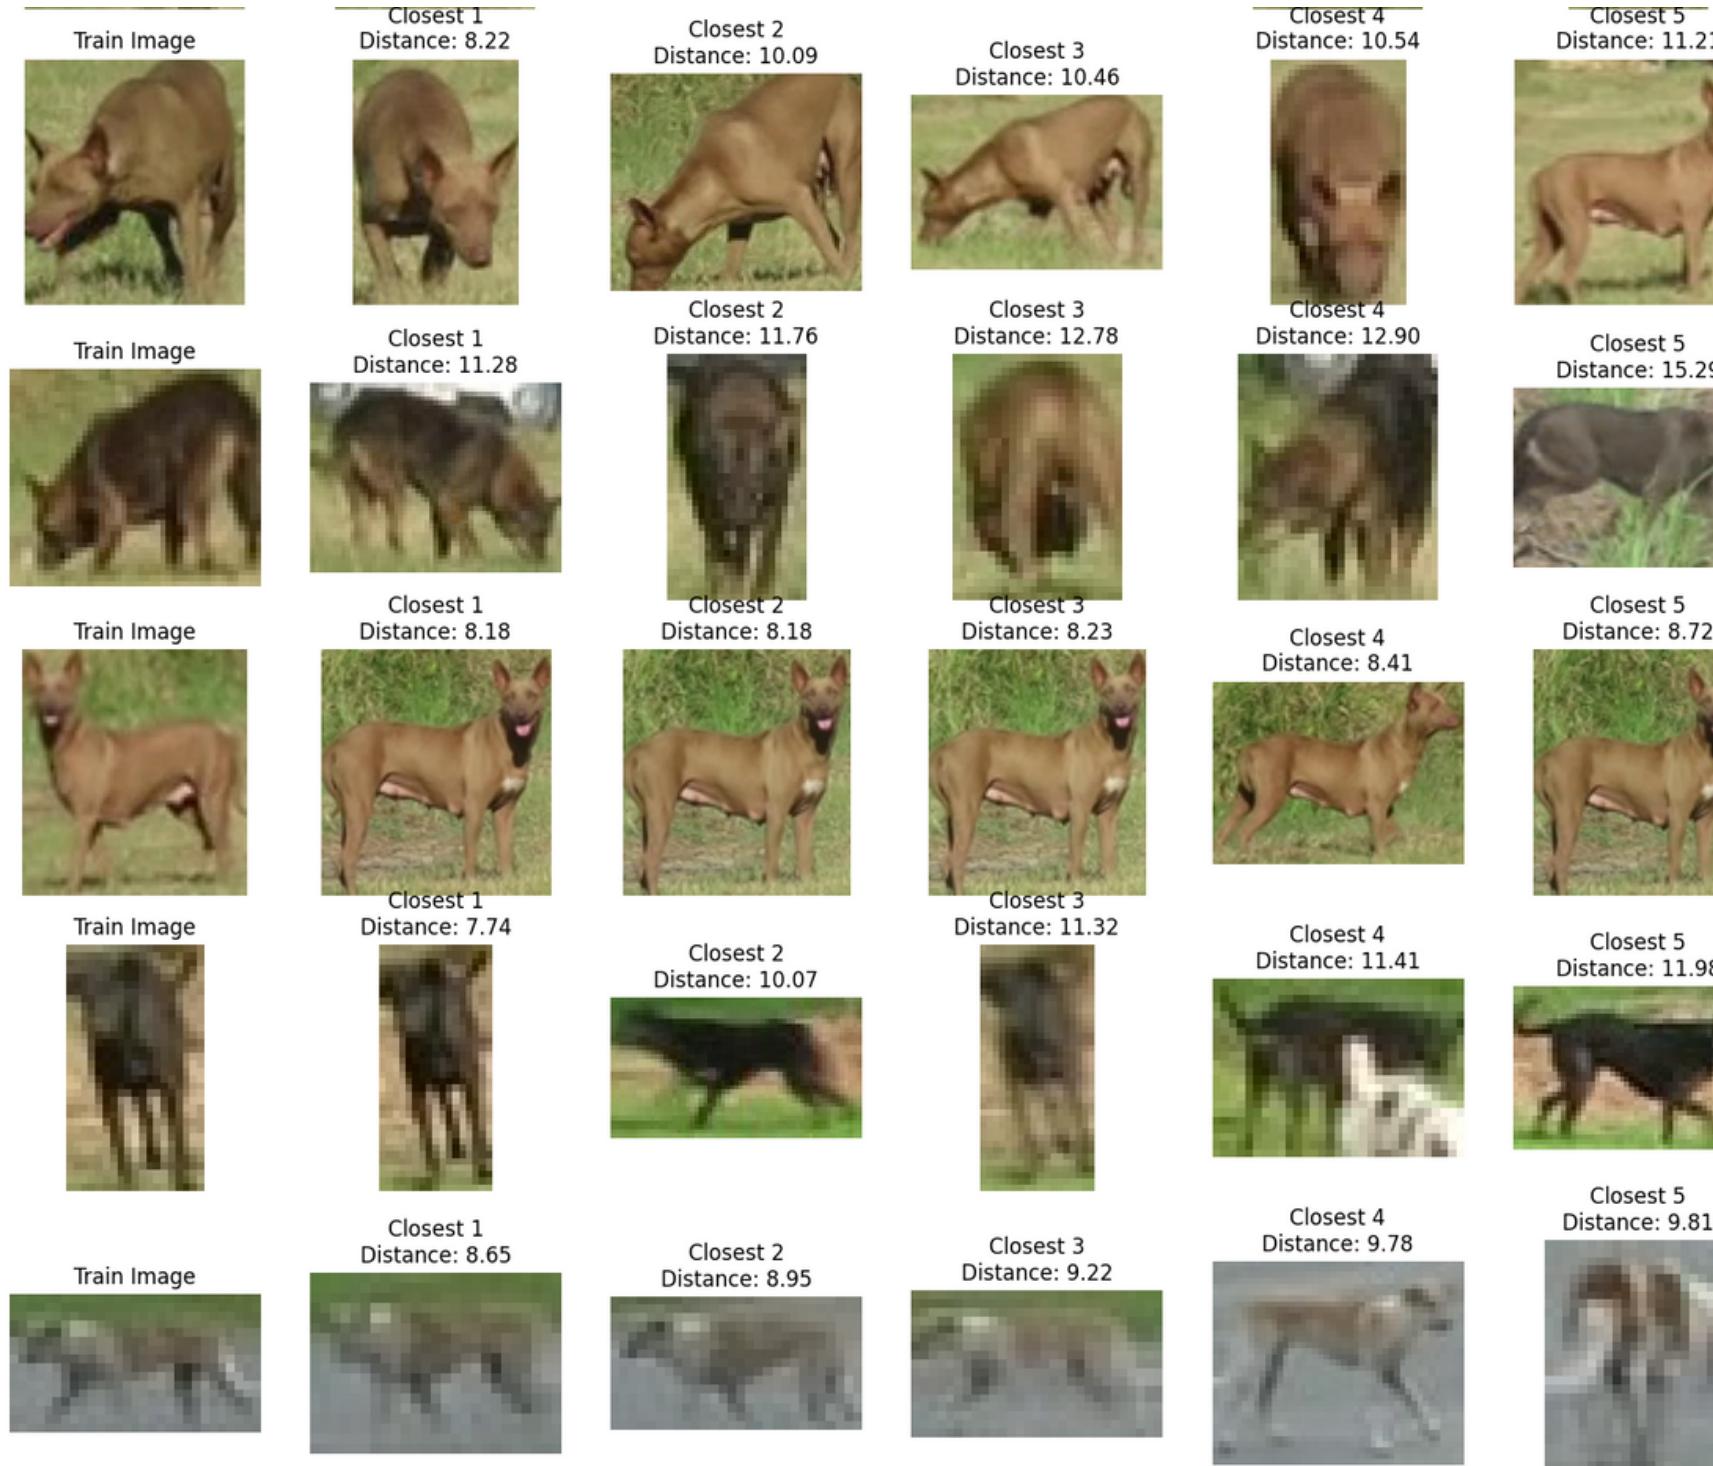


Cosine Similarity



Euclidean distance

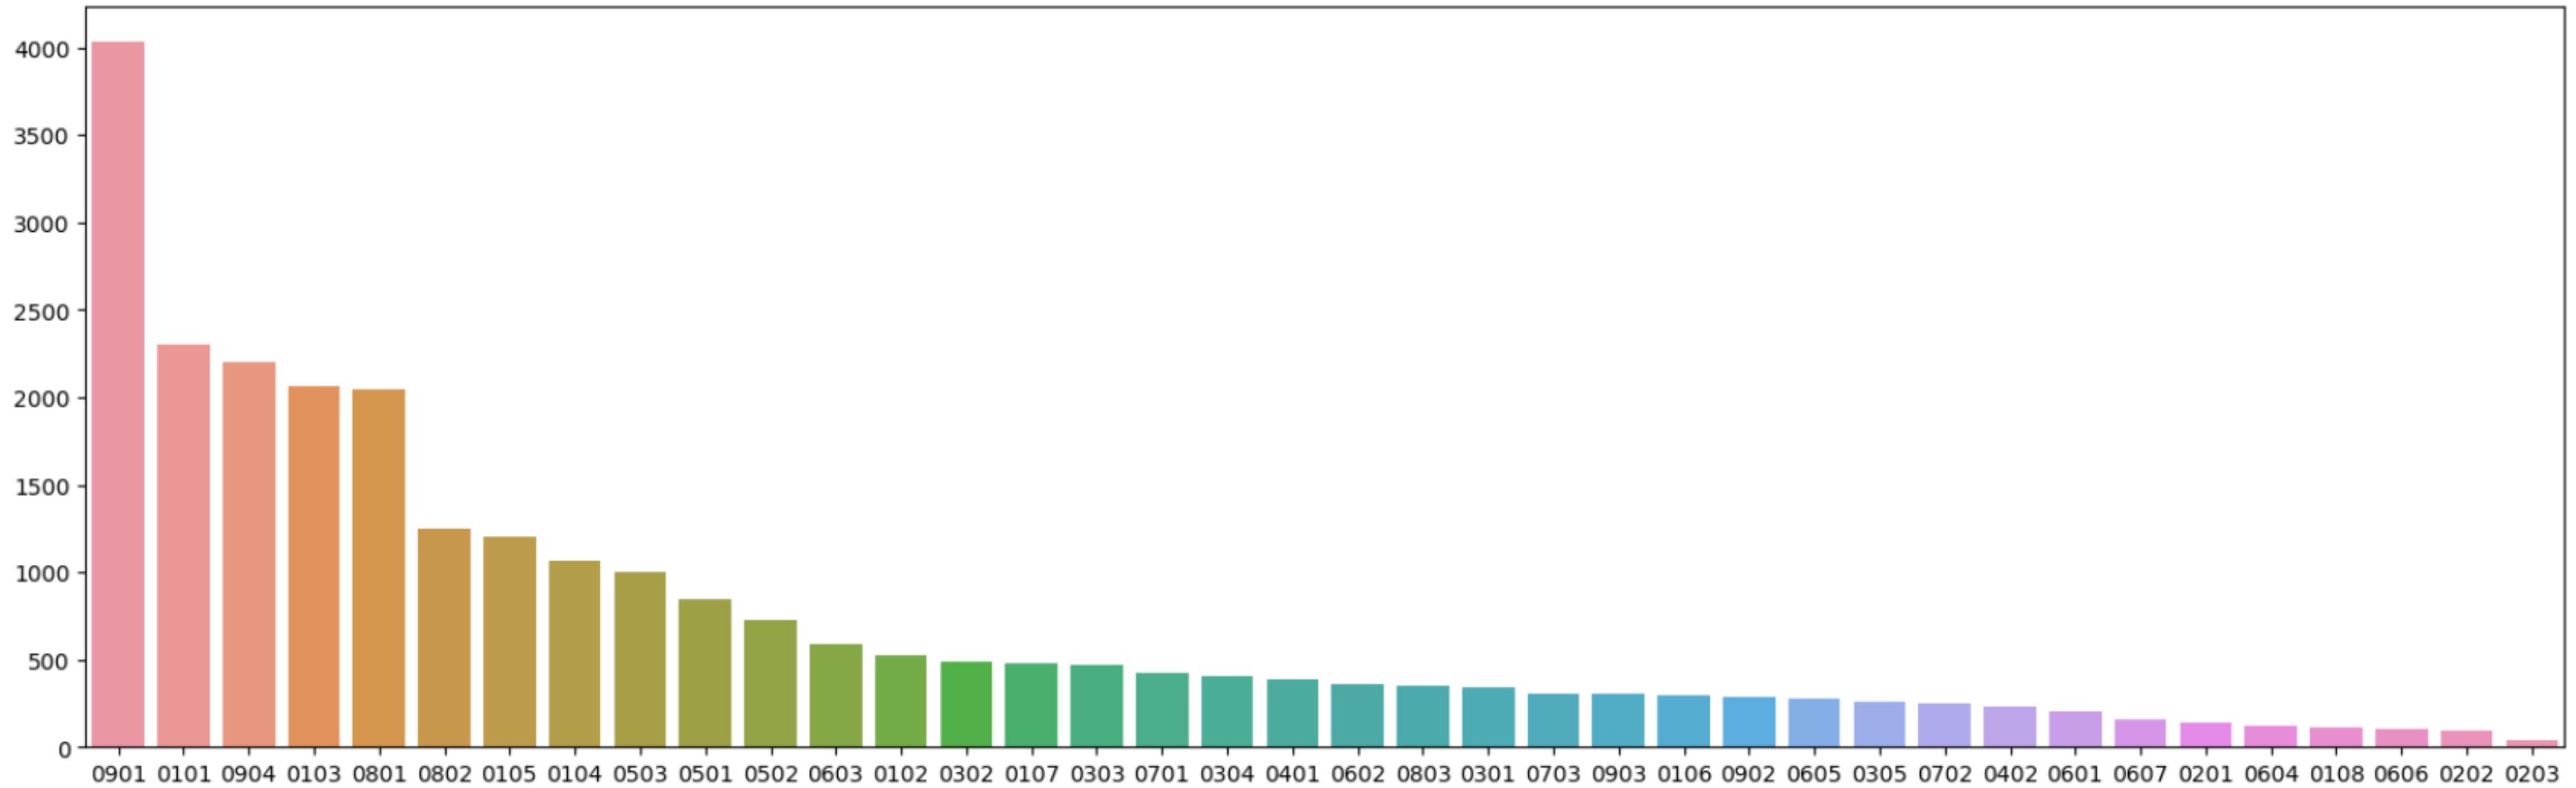
0.98 accuracy with test dataset



Generalization

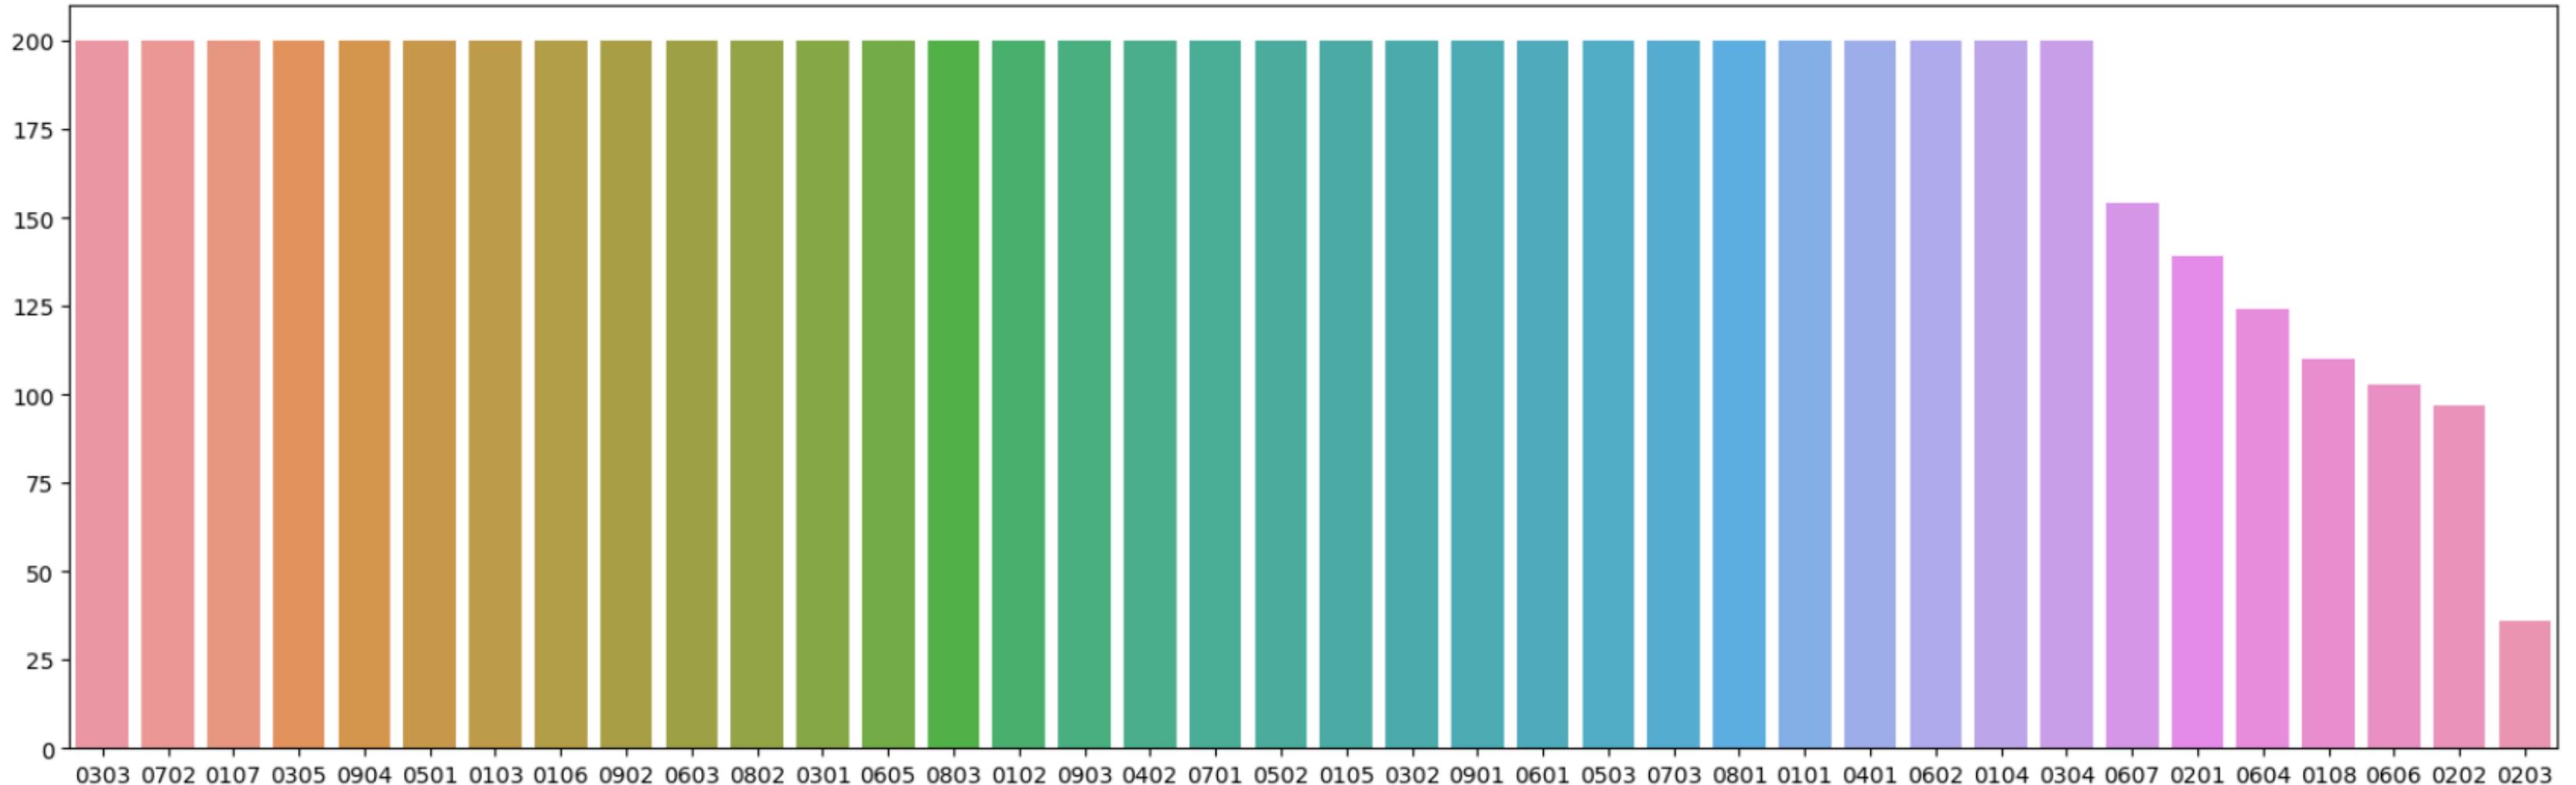
Can it generalized well with small training set?

Data: train



Total images = 26737 class = 39

Data: train

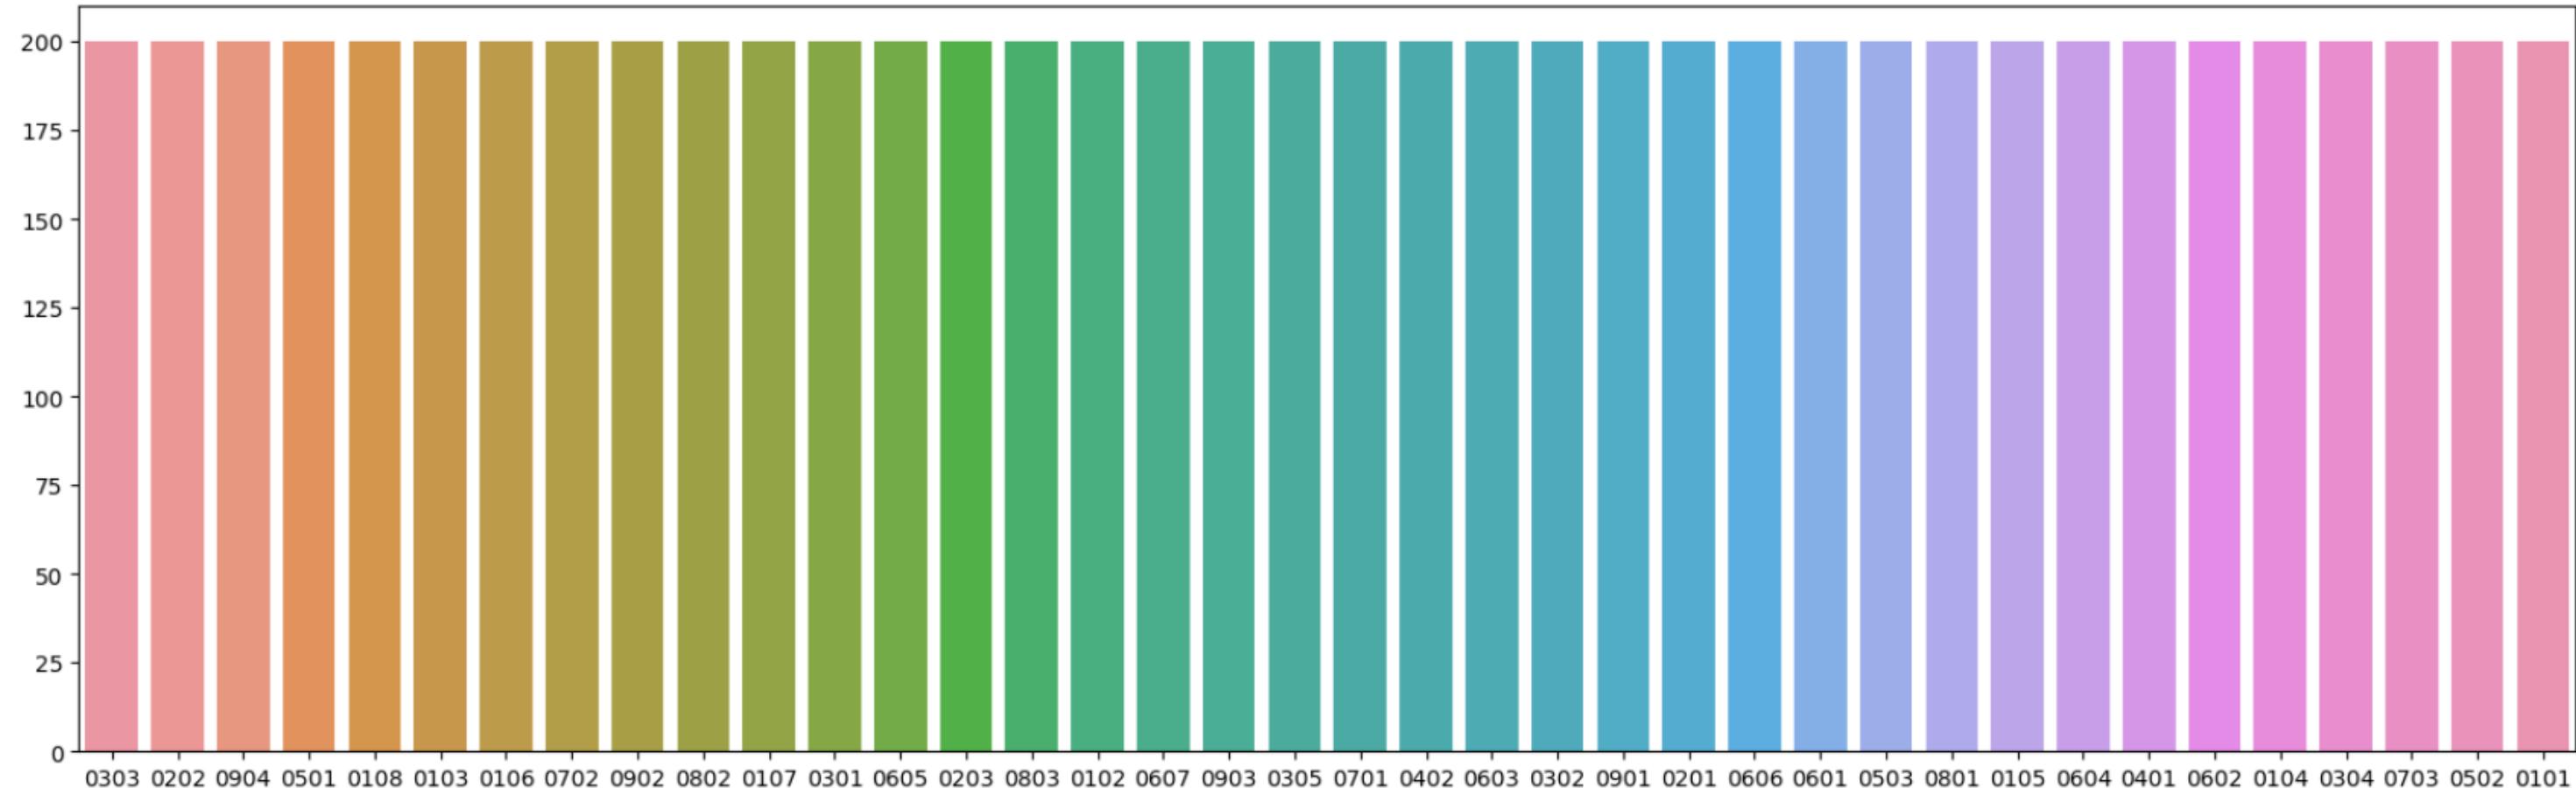


```

def get_largest_n_images(df, n=Tagert_sample):
    result_df = pd.DataFrame()
    for dog_id in df['DogID'].unique():
        dog_df = df[df['DogID'] == dog_id]
        largest_images = dog_df.nlargest(n, 'W' if dog_df['W'].max() >= dog_df['H'].max() else 'H')
        result_df = result_df.append(largest_images)
    return result_df
  
```

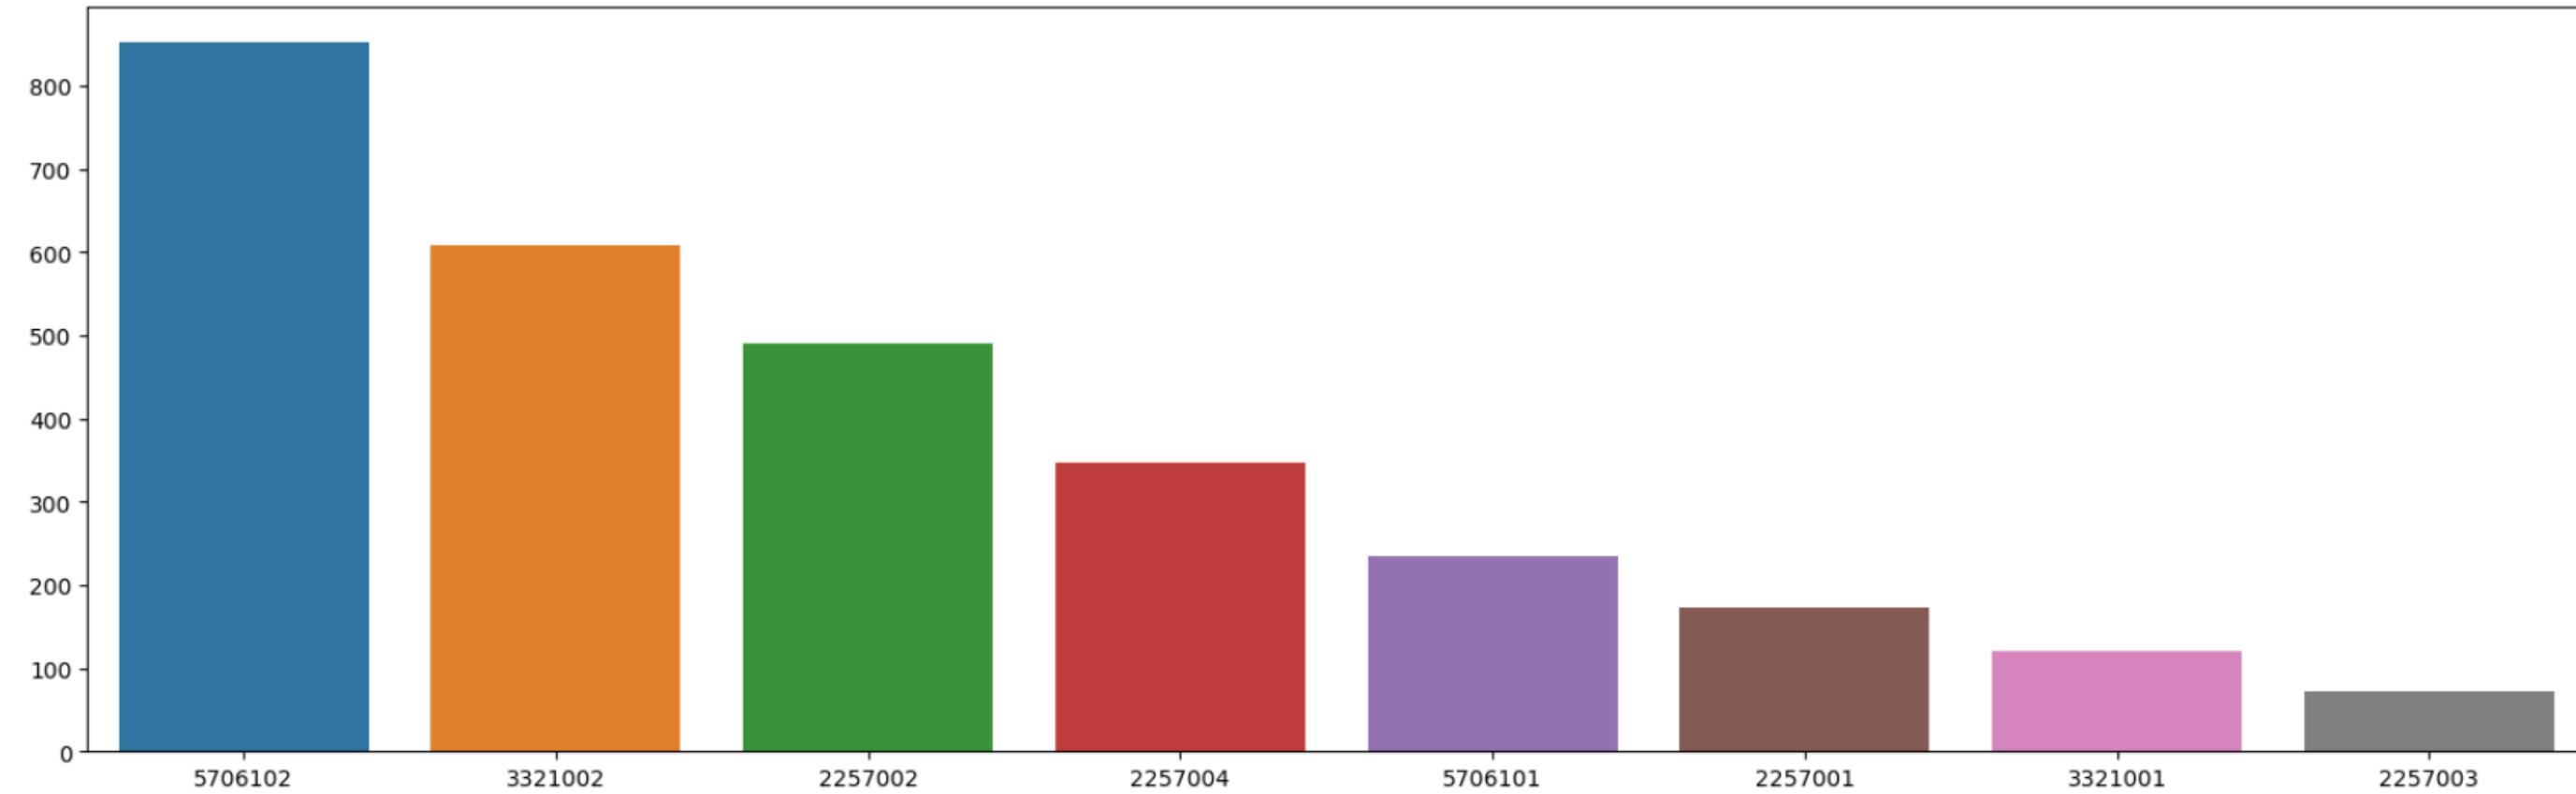
Data: train

Total images = 39*200 class = 39

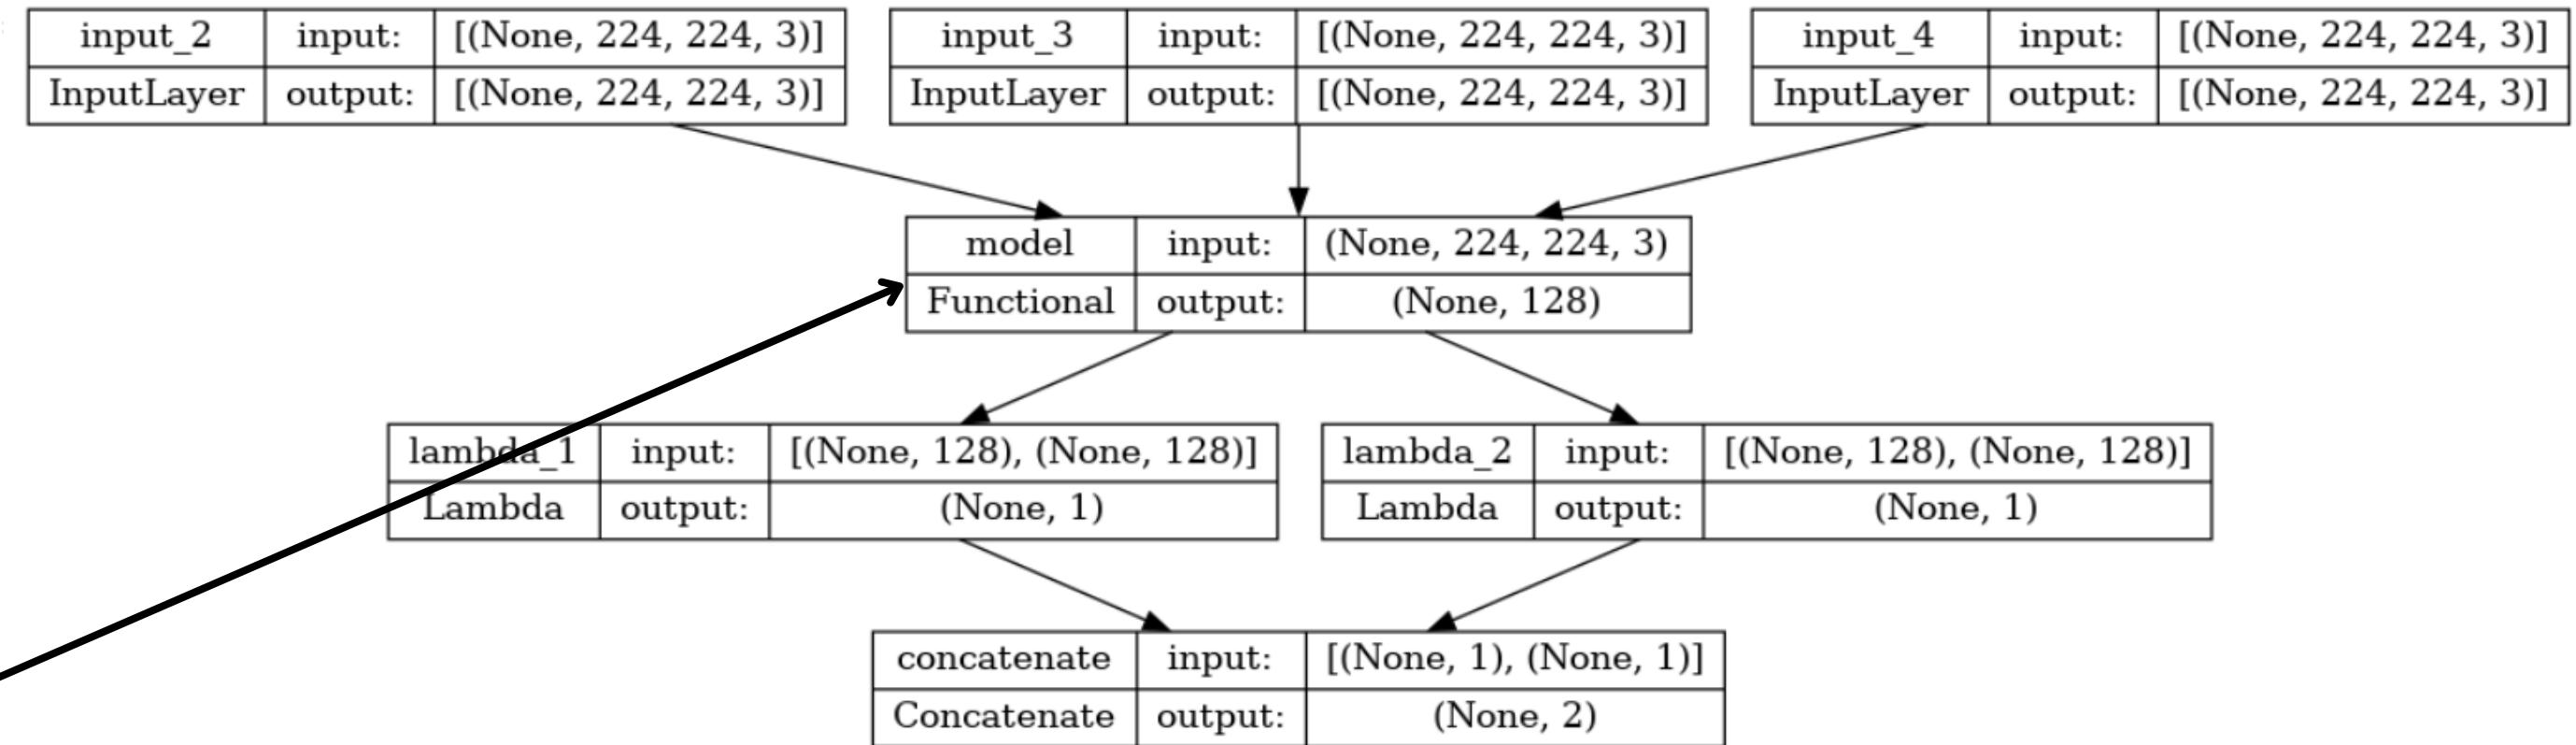


```
data_augmentation = tf.keras.Sequential([
    layers.experimental.preprocessing.Rescaling(1./255),
    tf.keras.layers.RandomFlip('horizontal'),
    tf.keras.layers.RandomRotation(0.1),
    tf.keras.layers.RandomTranslation(height_factor=0.1, width_factor=0.1, fill_mode="reflect"),
    tf.keras.layers.RandomContrast(factor=0.2),
    tf.keras.layers.RandomZoom(height_factor=(-0.1, 0.1), width_factor=(-0.1, 0.1)),
])
```

Data: test unseen



My Siamese Network (inceptionResnetv2)



```

def create_embedding_model():
    base_cnn = applications.inception_resnet_v2.InceptionResNetV2(
        weights='imagenet',
        include_top=False,
        input_shape=(*IMG_SIZE, 3)
    )

    x = layers.GlobalAveragePooling2D()(base_cnn.output)
    x = layers.Dense(1024, activation='relu')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Dense(512, activation='relu')(x)
    x = layers.BatchNormalization()(x)
    x = layers.Dense(256, activation='relu')(x)
    x = layers.BatchNormalization()(x)
    output = layers.Dense(128)(x)
    output = layers.Lambda(lambda x: K.l2_normalize(x, axis=-1))(output) # Add L2 normalization

    return Model(inputs=base_cnn.input, outputs=output)
  
```

```

def create_siamese_model(embedding_model):
    anchor_input = layers.Input(shape=(*IMG_SIZE, 3))
    positive_input = layers.Input(shape=(*IMG_SIZE, 3))
    negative_input = layers.Input(shape=(*IMG_SIZE, 3))

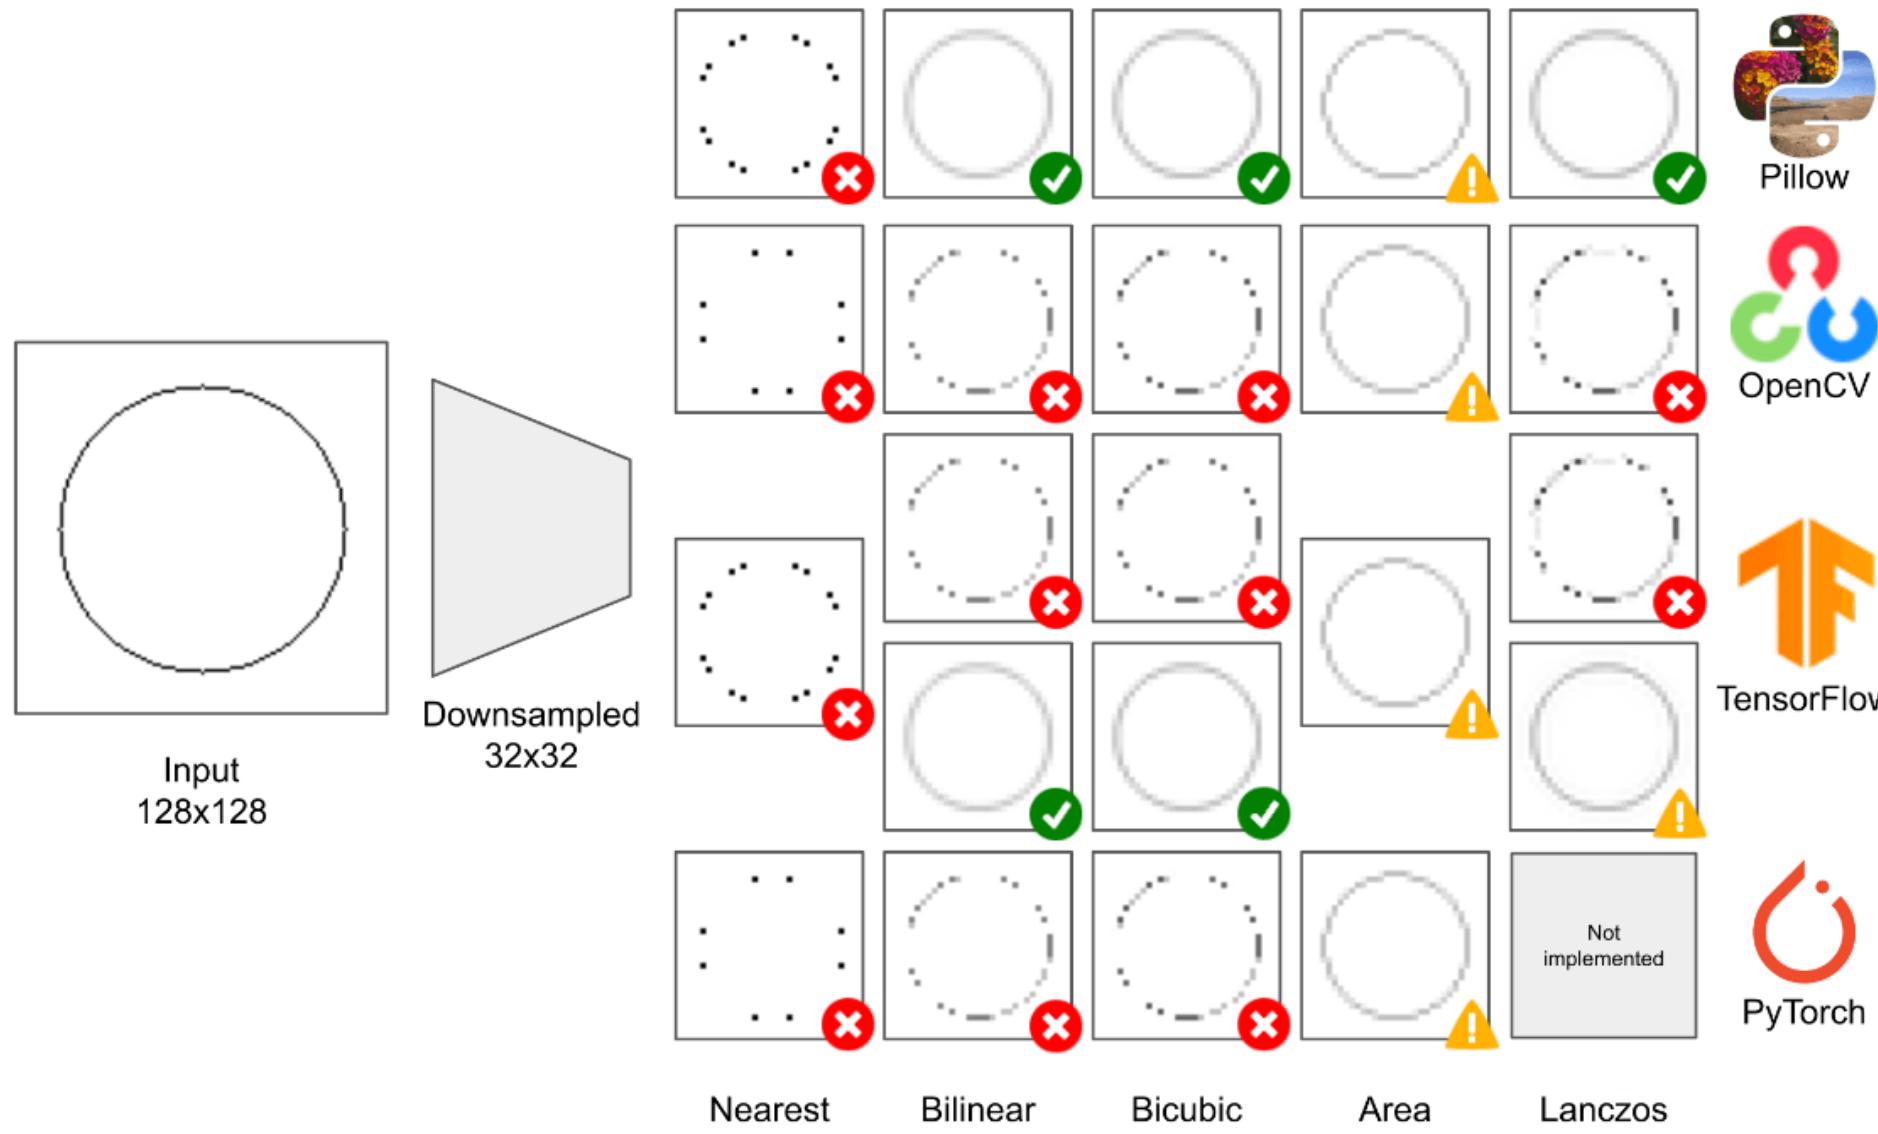
    anchor_embedding = embedding_model(anchor_input)
    positive_embedding = embedding_model(positive_input)
    negative_embedding = embedding_model(negative_input)

    positive_distance = layers.Lambda(euclidean_distance)([anchor_embedding, positive_embedding])
    negative_distance = layers.Lambda(euclidean_distance)([anchor_embedding, negative_embedding])

    stacked_dists = layers.Concatenate()([positive_distance, negative_distance])
    siamese_model = Model([anchor_input, positive_input, negative_input], stacked_dists)

    siamese_model.compile(optimizer=tf.keras.optimizers.Adam(1e-5), loss=triplet_loss, metrics=[triplet_accuracy])
    return siamese_model
  
```

Data Generator Code Snippet



```

def load_image(image_path, img_size=IMG_SIZE):
    img = plt.imread(image_path)
    img = np.array(Image.fromarray(img).resize(img_size, resample=Image.BILINEAR))
    return img / 255.0

def data_generator(data_df, batch_size=BATCH_SIZE):
    while True:
        data_df = data_df.sample(frac=1).reset_index(drop=True) # Shuffle the dataframe
        for i in range(0, len(data_df), batch_size):
            batch_df = data_df[i:i + batch_size]
            anchor_batch, positive_batch, negative_batch = [], [], []
            for _, row in batch_df.iterrows():
                anchor_img_path = row['filepath']
                anchor_dog_id = row['DogID']

                # Get positive image
                positive_df = data_df[data_df['DogID'] == anchor_dog_id].sample(n=2)
                positive_row = positive_df.iloc[0] if positive_df.iloc[0]['filepath'] != anchor_img_path else positive_df.iloc[1]
                positive_img_path = positive_row['filepath']

                # Get negative image
                negative_df = data_df[data_df['DogID'] != anchor_dog_id].sample(1)
                negative_img_path = negative_df.iloc[0]['filepath']

                anchor_batch.append(load_image(anchor_img_path))
                positive_batch.append(load_image(positive_img_path))
                negative_batch.append(load_image(negative_img_path))

            anchor_batch = np.array(anchor_batch)
            positive_batch = np.array(positive_batch)
            negative_batch = np.array(negative_batch)
            dummy_labels = np.zeros((batch_size, 1)) # Add dummy labels
            yield [anchor_batch, positive_batch, negative_batch], dummy_labels
  
```

My Siamese Network (inceptionResnetv2)

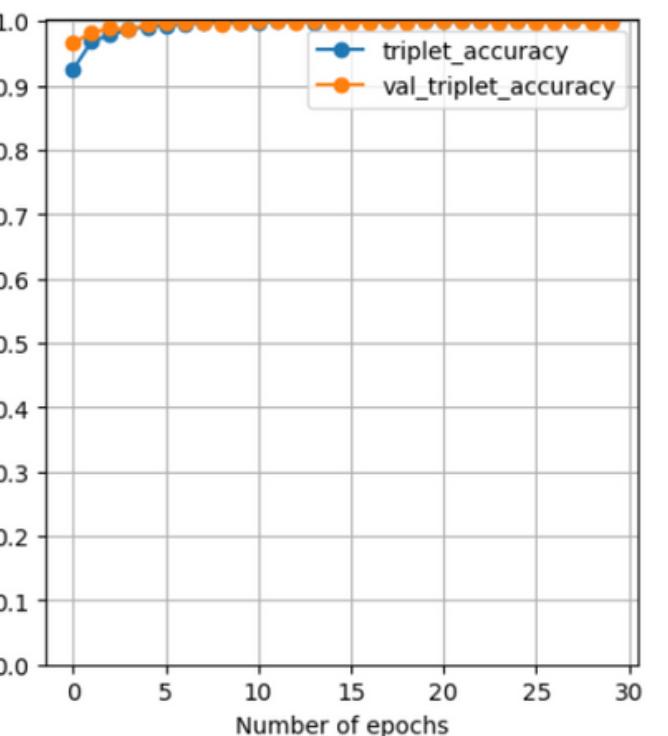
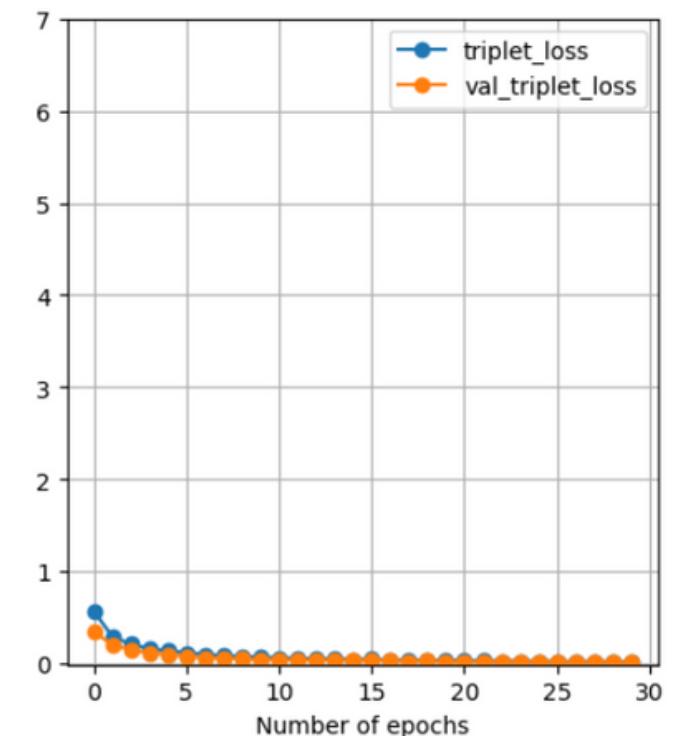
```
print(f'This is the number of trainable weights {len(embedding_model.trainable_weights)}')
```

This is the number of trainable weights 502

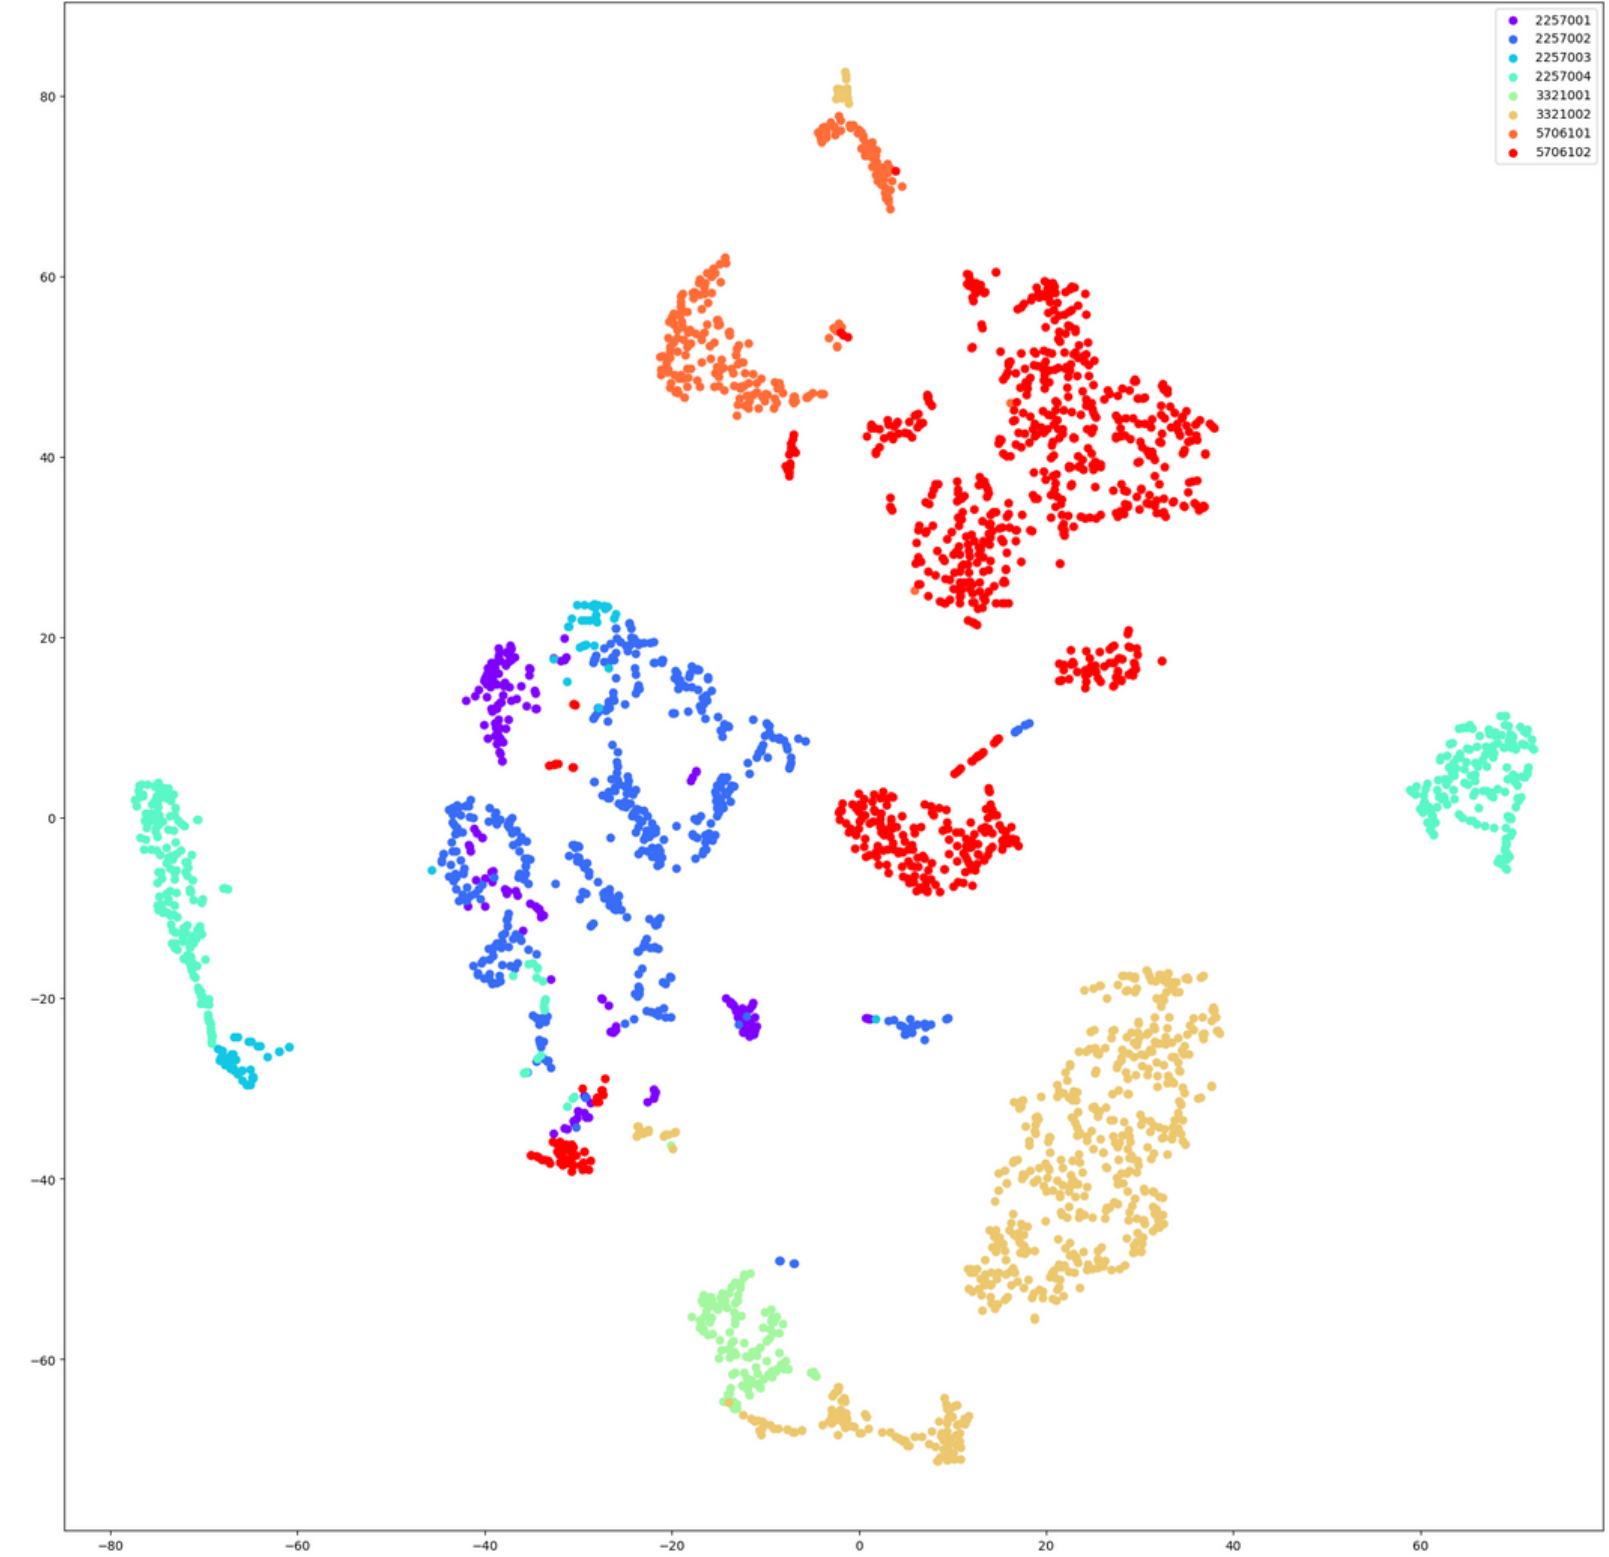
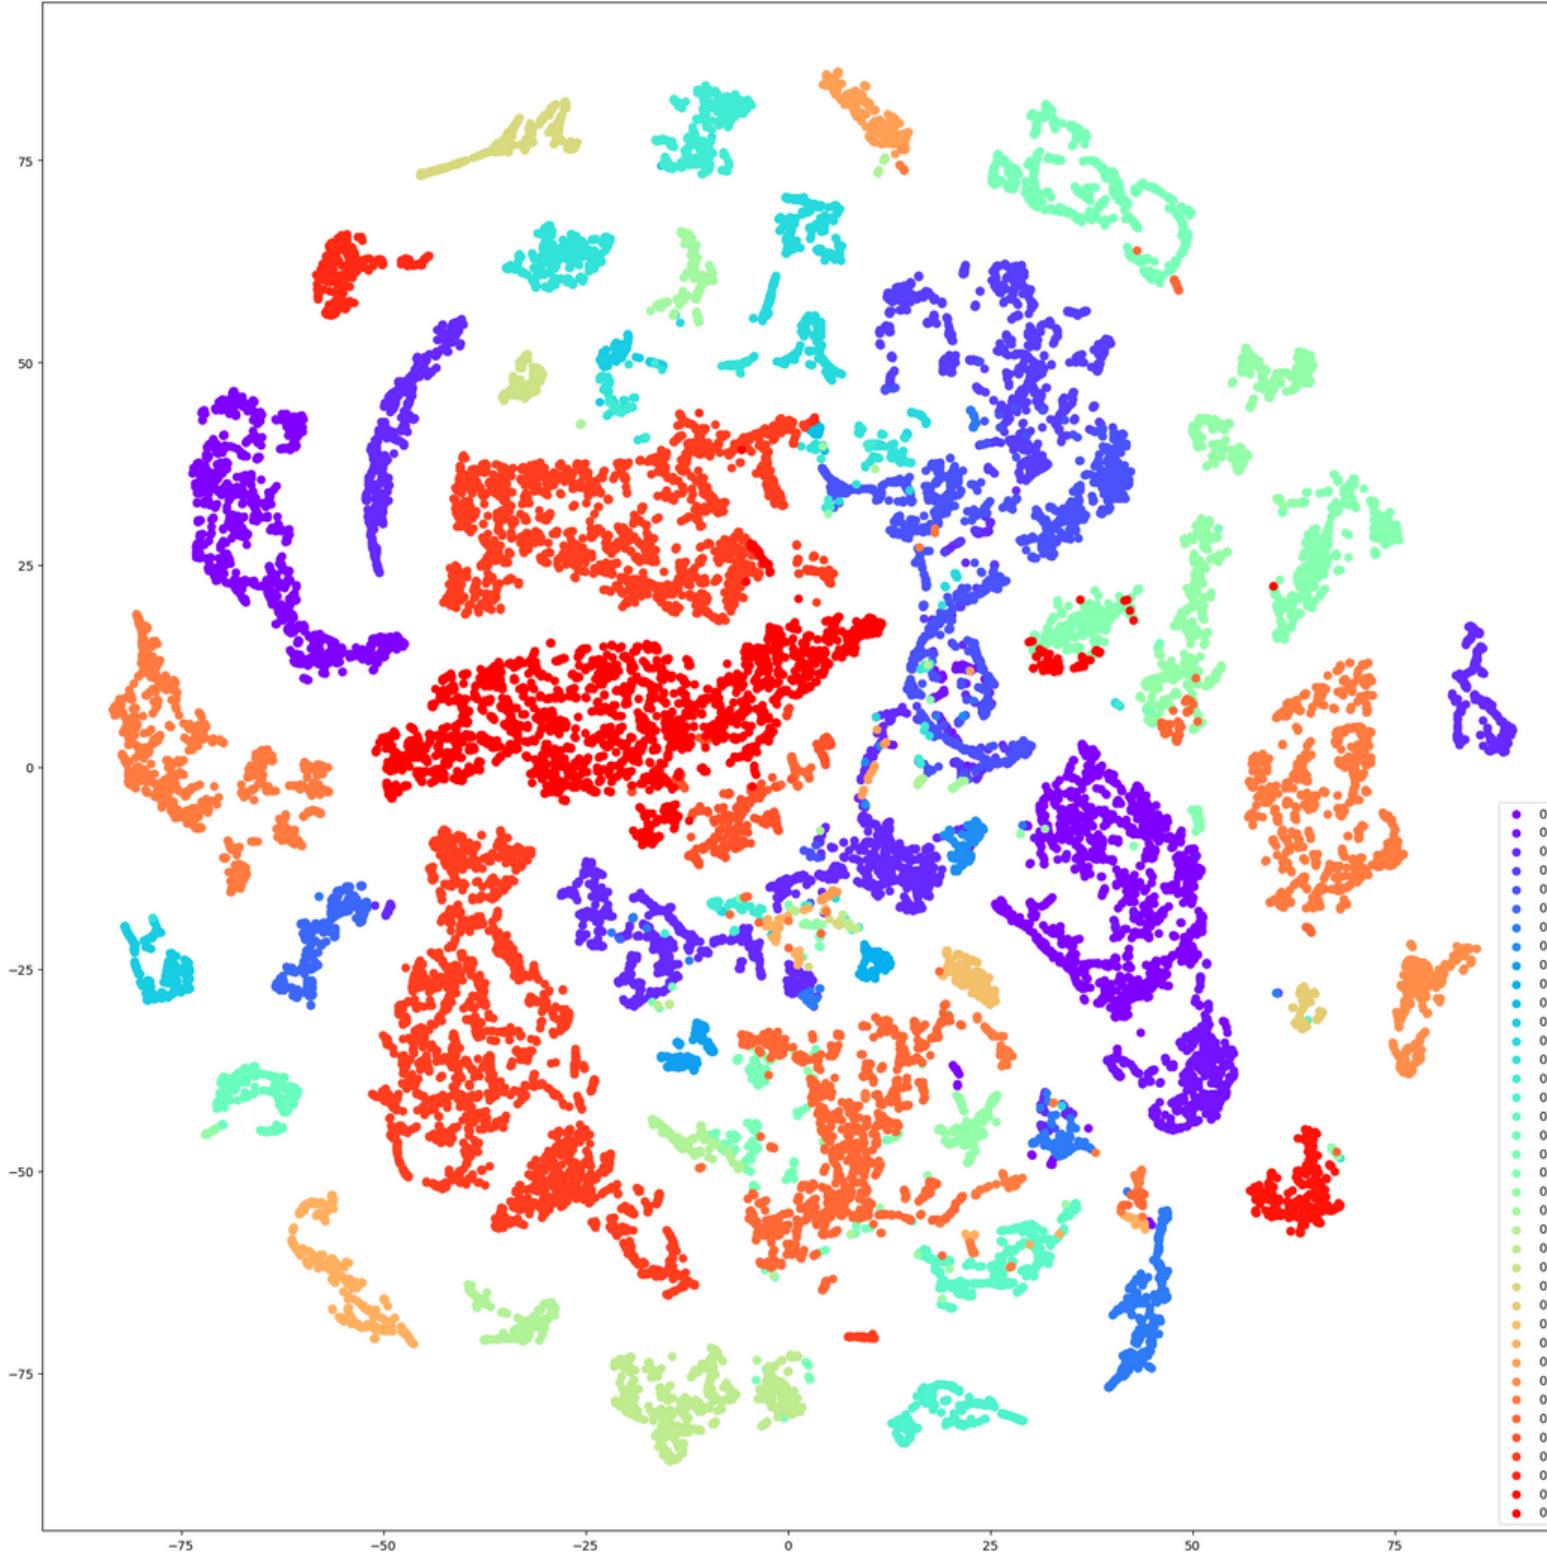
```
from tensorflow.keras.callbacks import EarlyStopping

# Define your early stopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

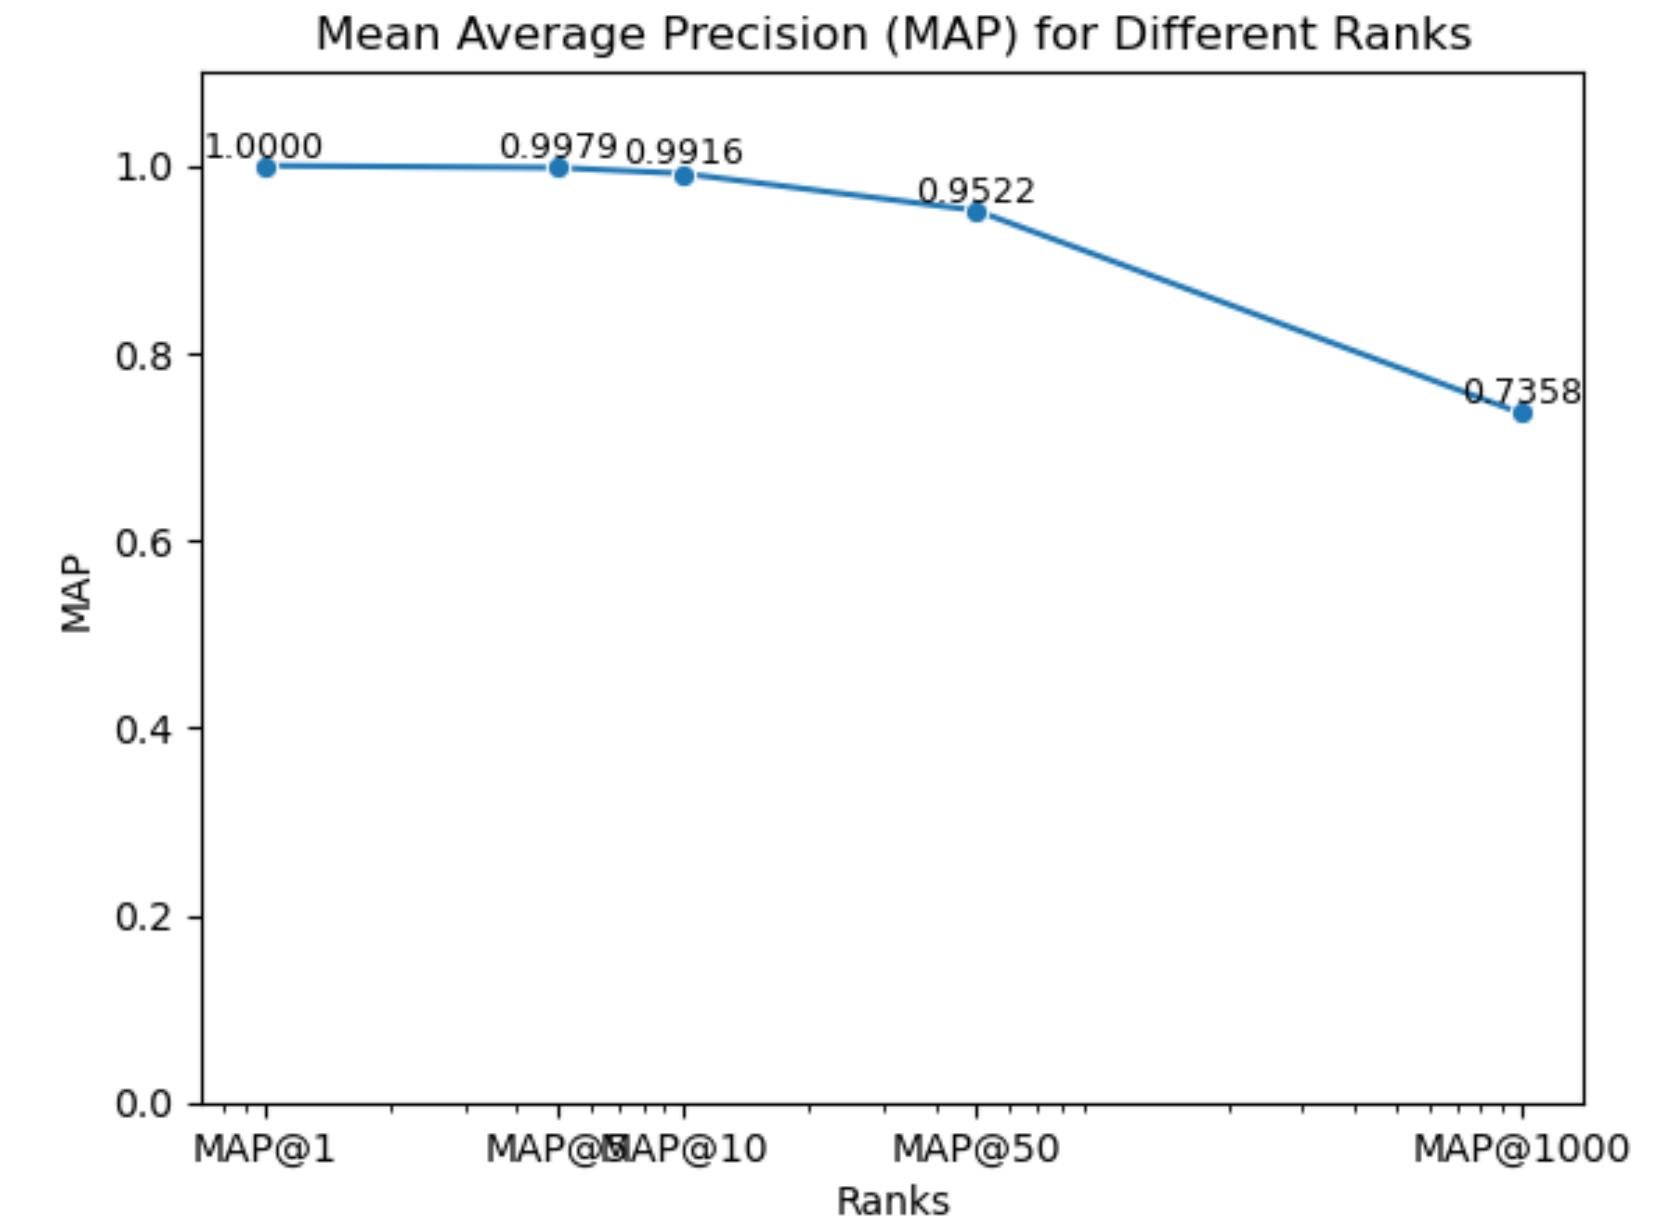
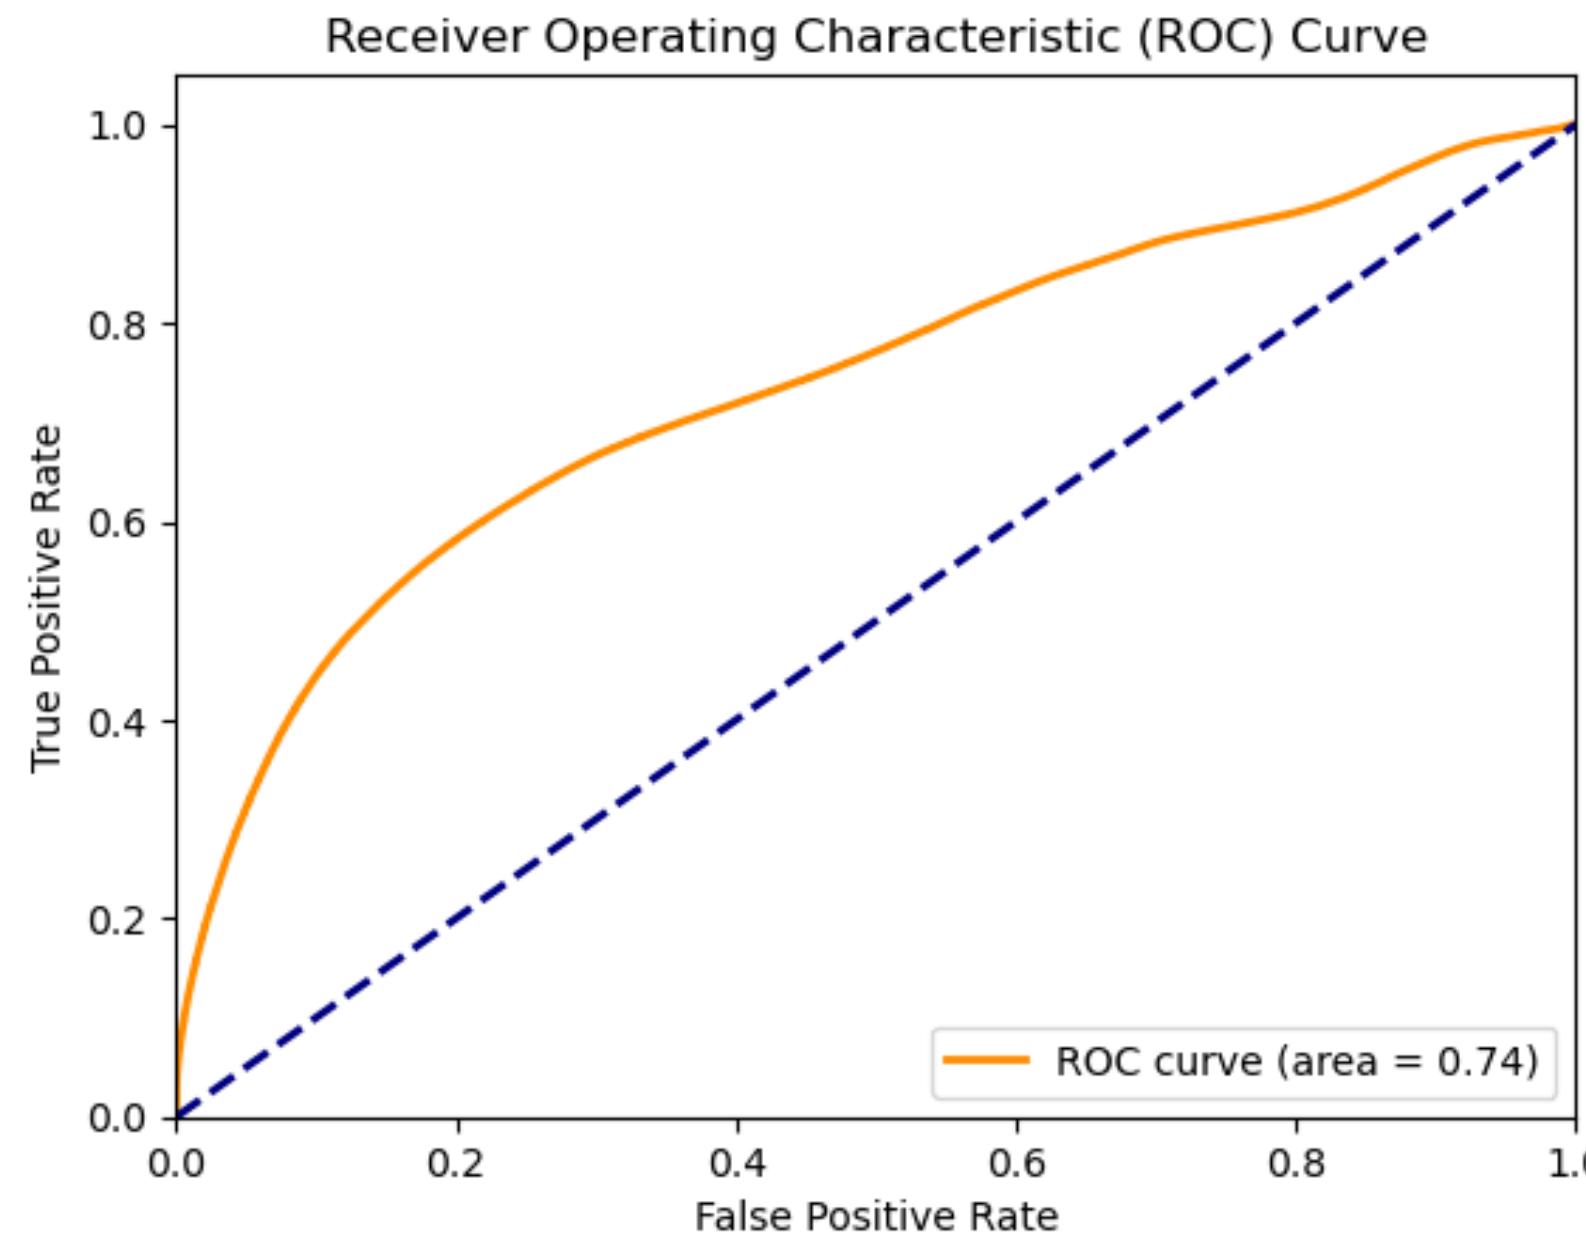
# Train the Siamese model with the learning rate scheduler and early stopping
history = siamese_model.fit(
    train_gen,
    steps_per_epoch=len(train_df) // BATCH_SIZE,
    validation_data=val_gen,
    validation_steps=len(val_df) // BATCH_SIZE,
    epochs=50,
    callbacks=[early_stopping]
)
```



My Siamese Network (inceptionResnetv2)



My Siamese Network (inceptionResnetv2)



Adjust trainable parameter

My Siamese Network (inceptionResnetv2)

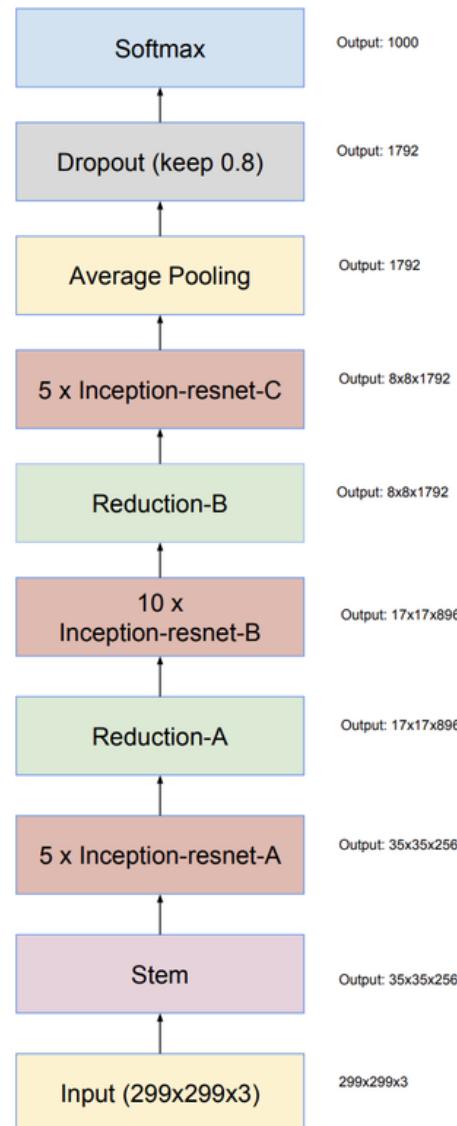


Figure 15. Schema for Inception-ResNet-v1 and Inception-ResNet-v2 networks. This schema applies to both networks but the underlying components differ. Inception-ResNet-v1 uses the blocks as described in Figures 14, 10, 7, 11, 12 and 13. Inception-ResNet-v2 uses the blocks as described in Figures 3, 16, 7, 17, 18 and 19. The output sizes in the diagram refer to the activation vector tensor shapes of Inception-ResNet-v1.

```
print(f'This is the number of trainable weights {len(embedding_model.trainable_weights)}')
```

This is the number of trainable weights 502

```
# Freeze all layers initially
for layer in embedding_model.layers:
    layer.trainable = False

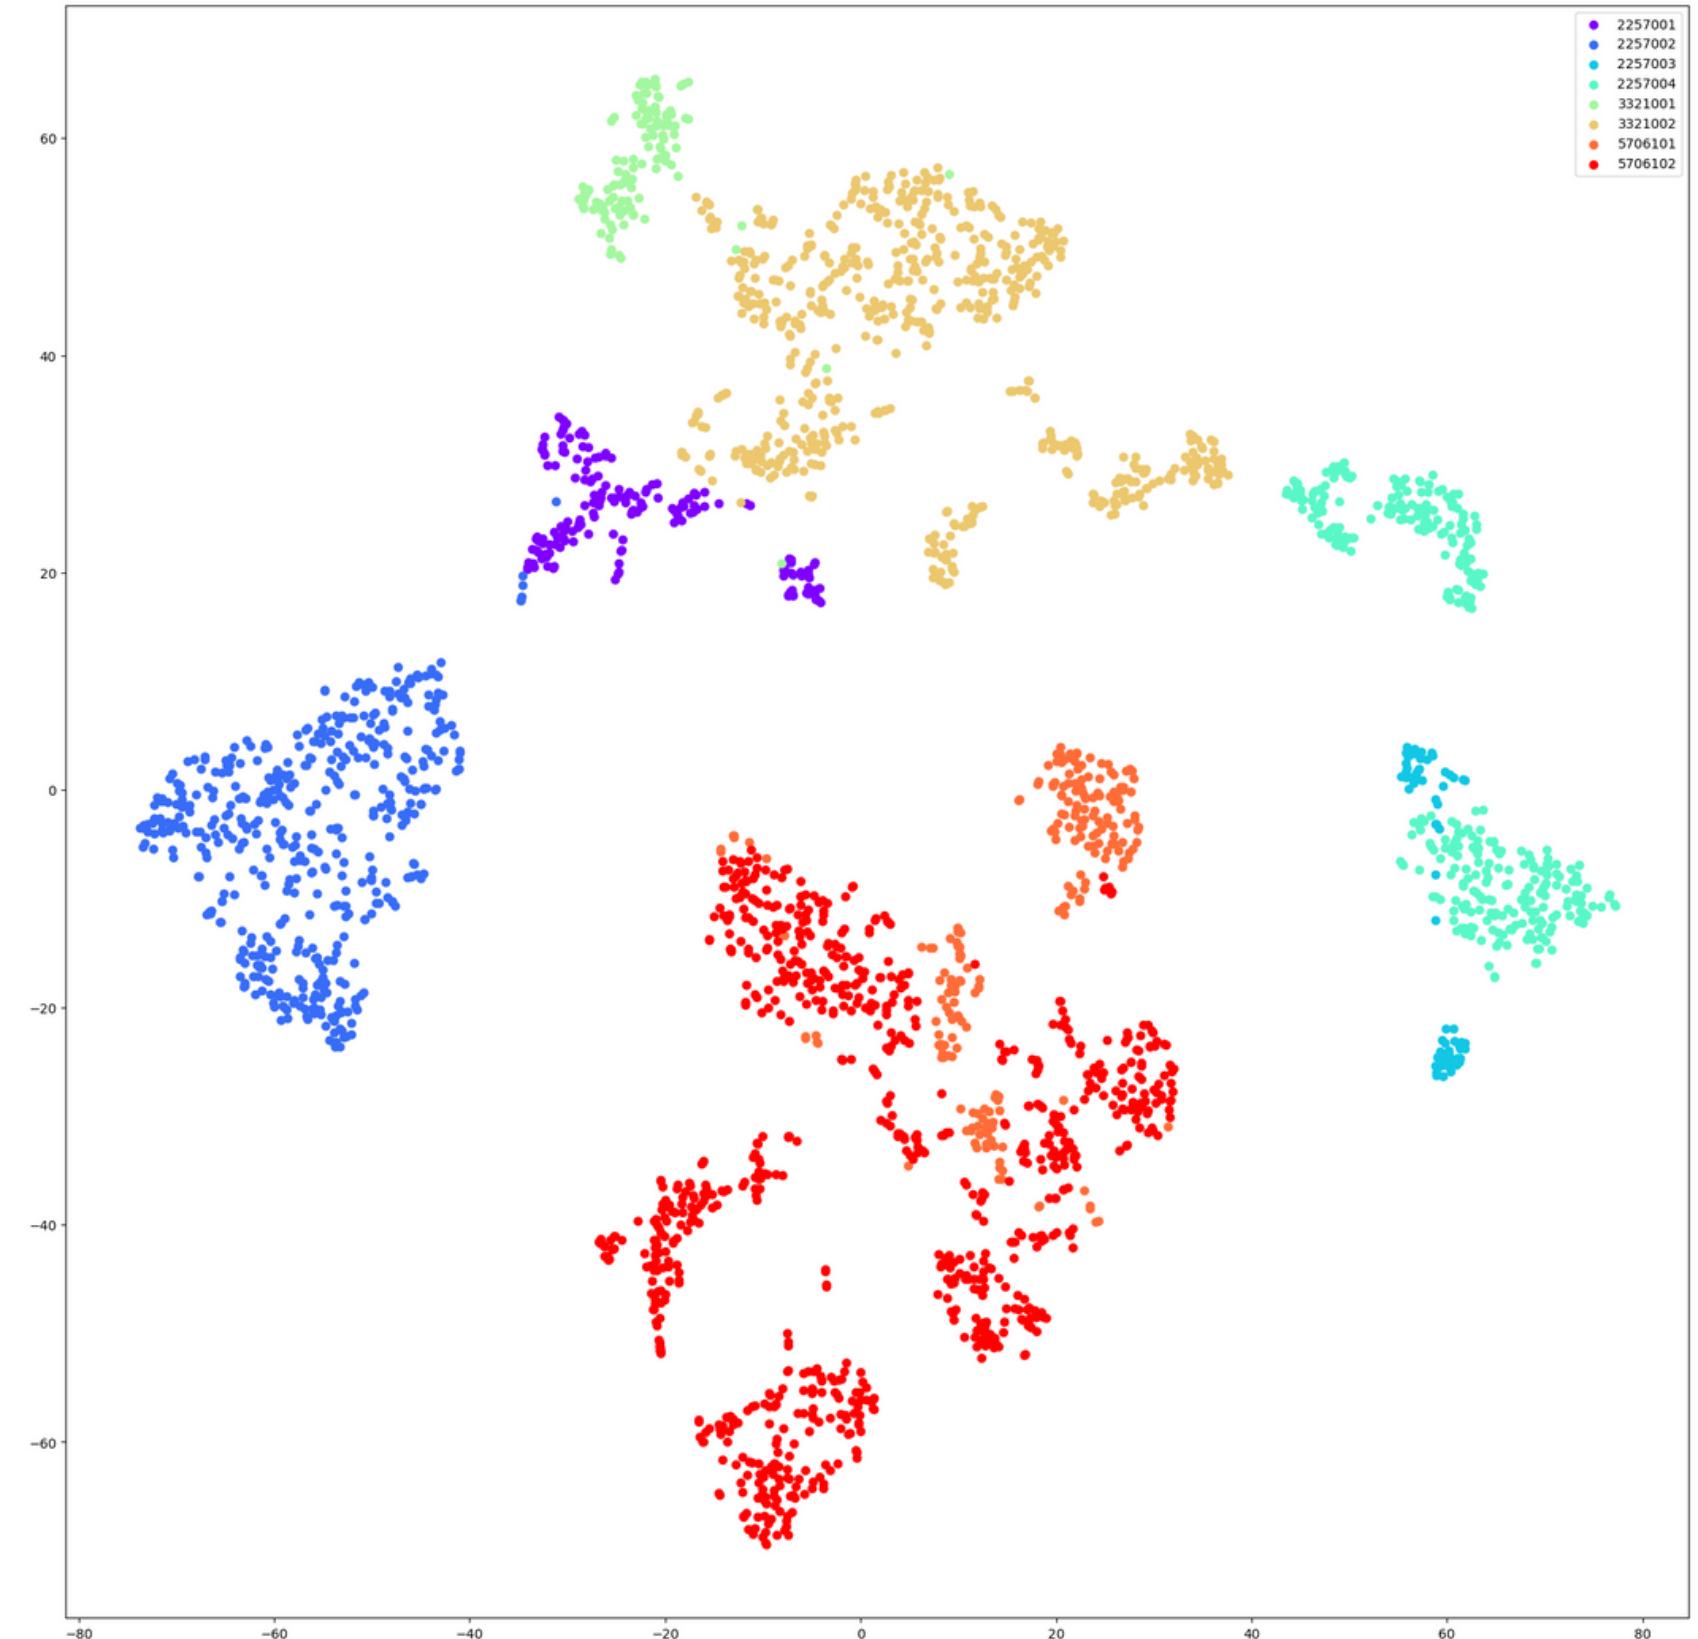
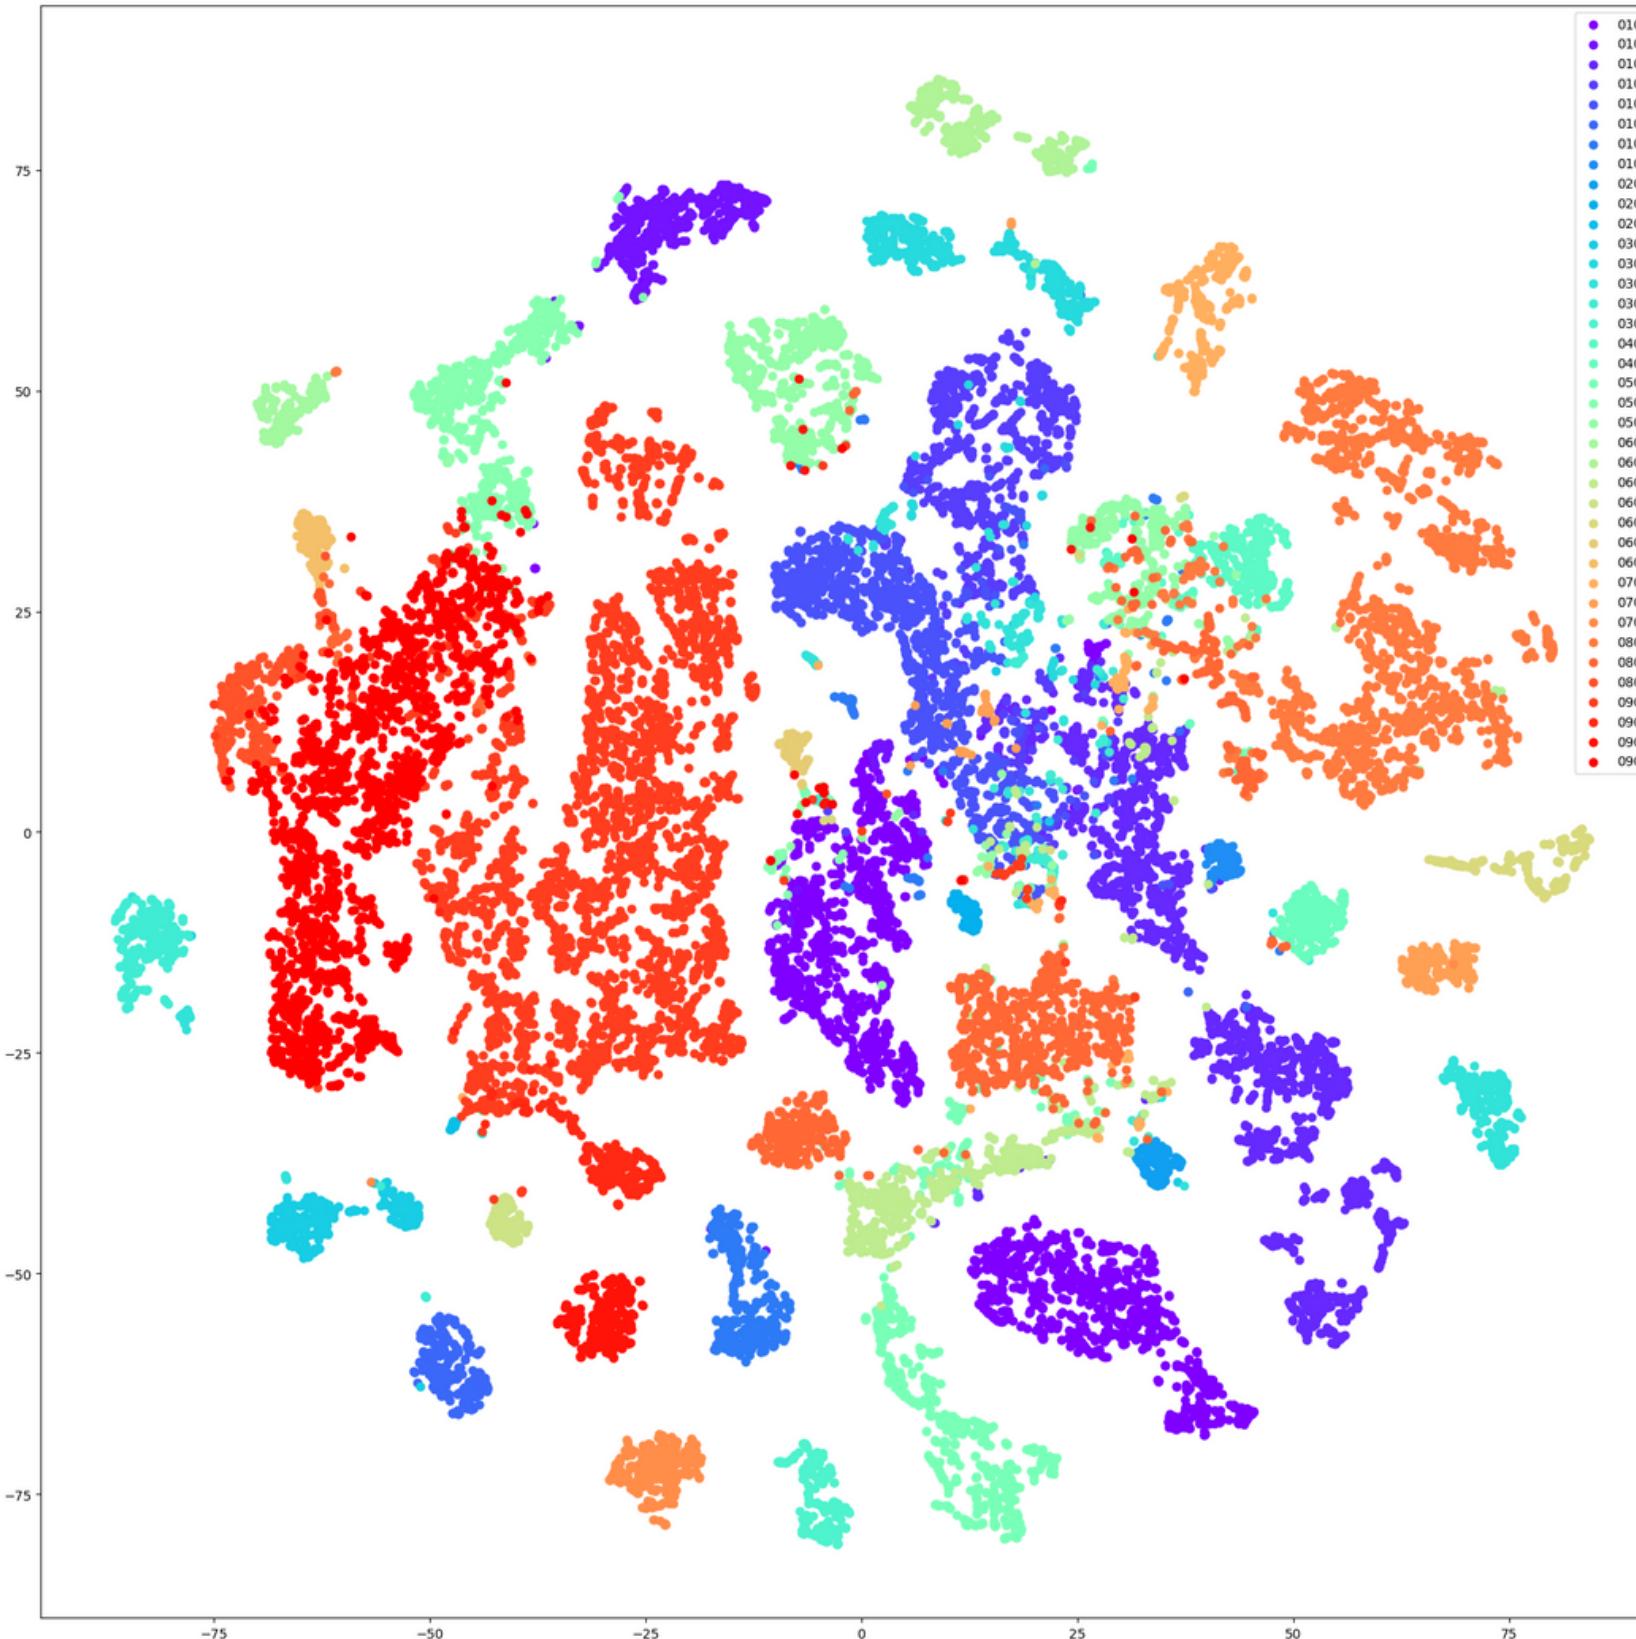
# Find the index of the first layer of the last Inception-C module
for i, layer in enumerate(embedding_model.layers):
    if layer.name == 'block8_1_conv':
        block8_start_index = i
        break

# Unfreeze the Block8 (last Inception-C module) and dense layers
for layer in embedding_model.layers[block8_start_index:]:
    layer.trainable = True
```

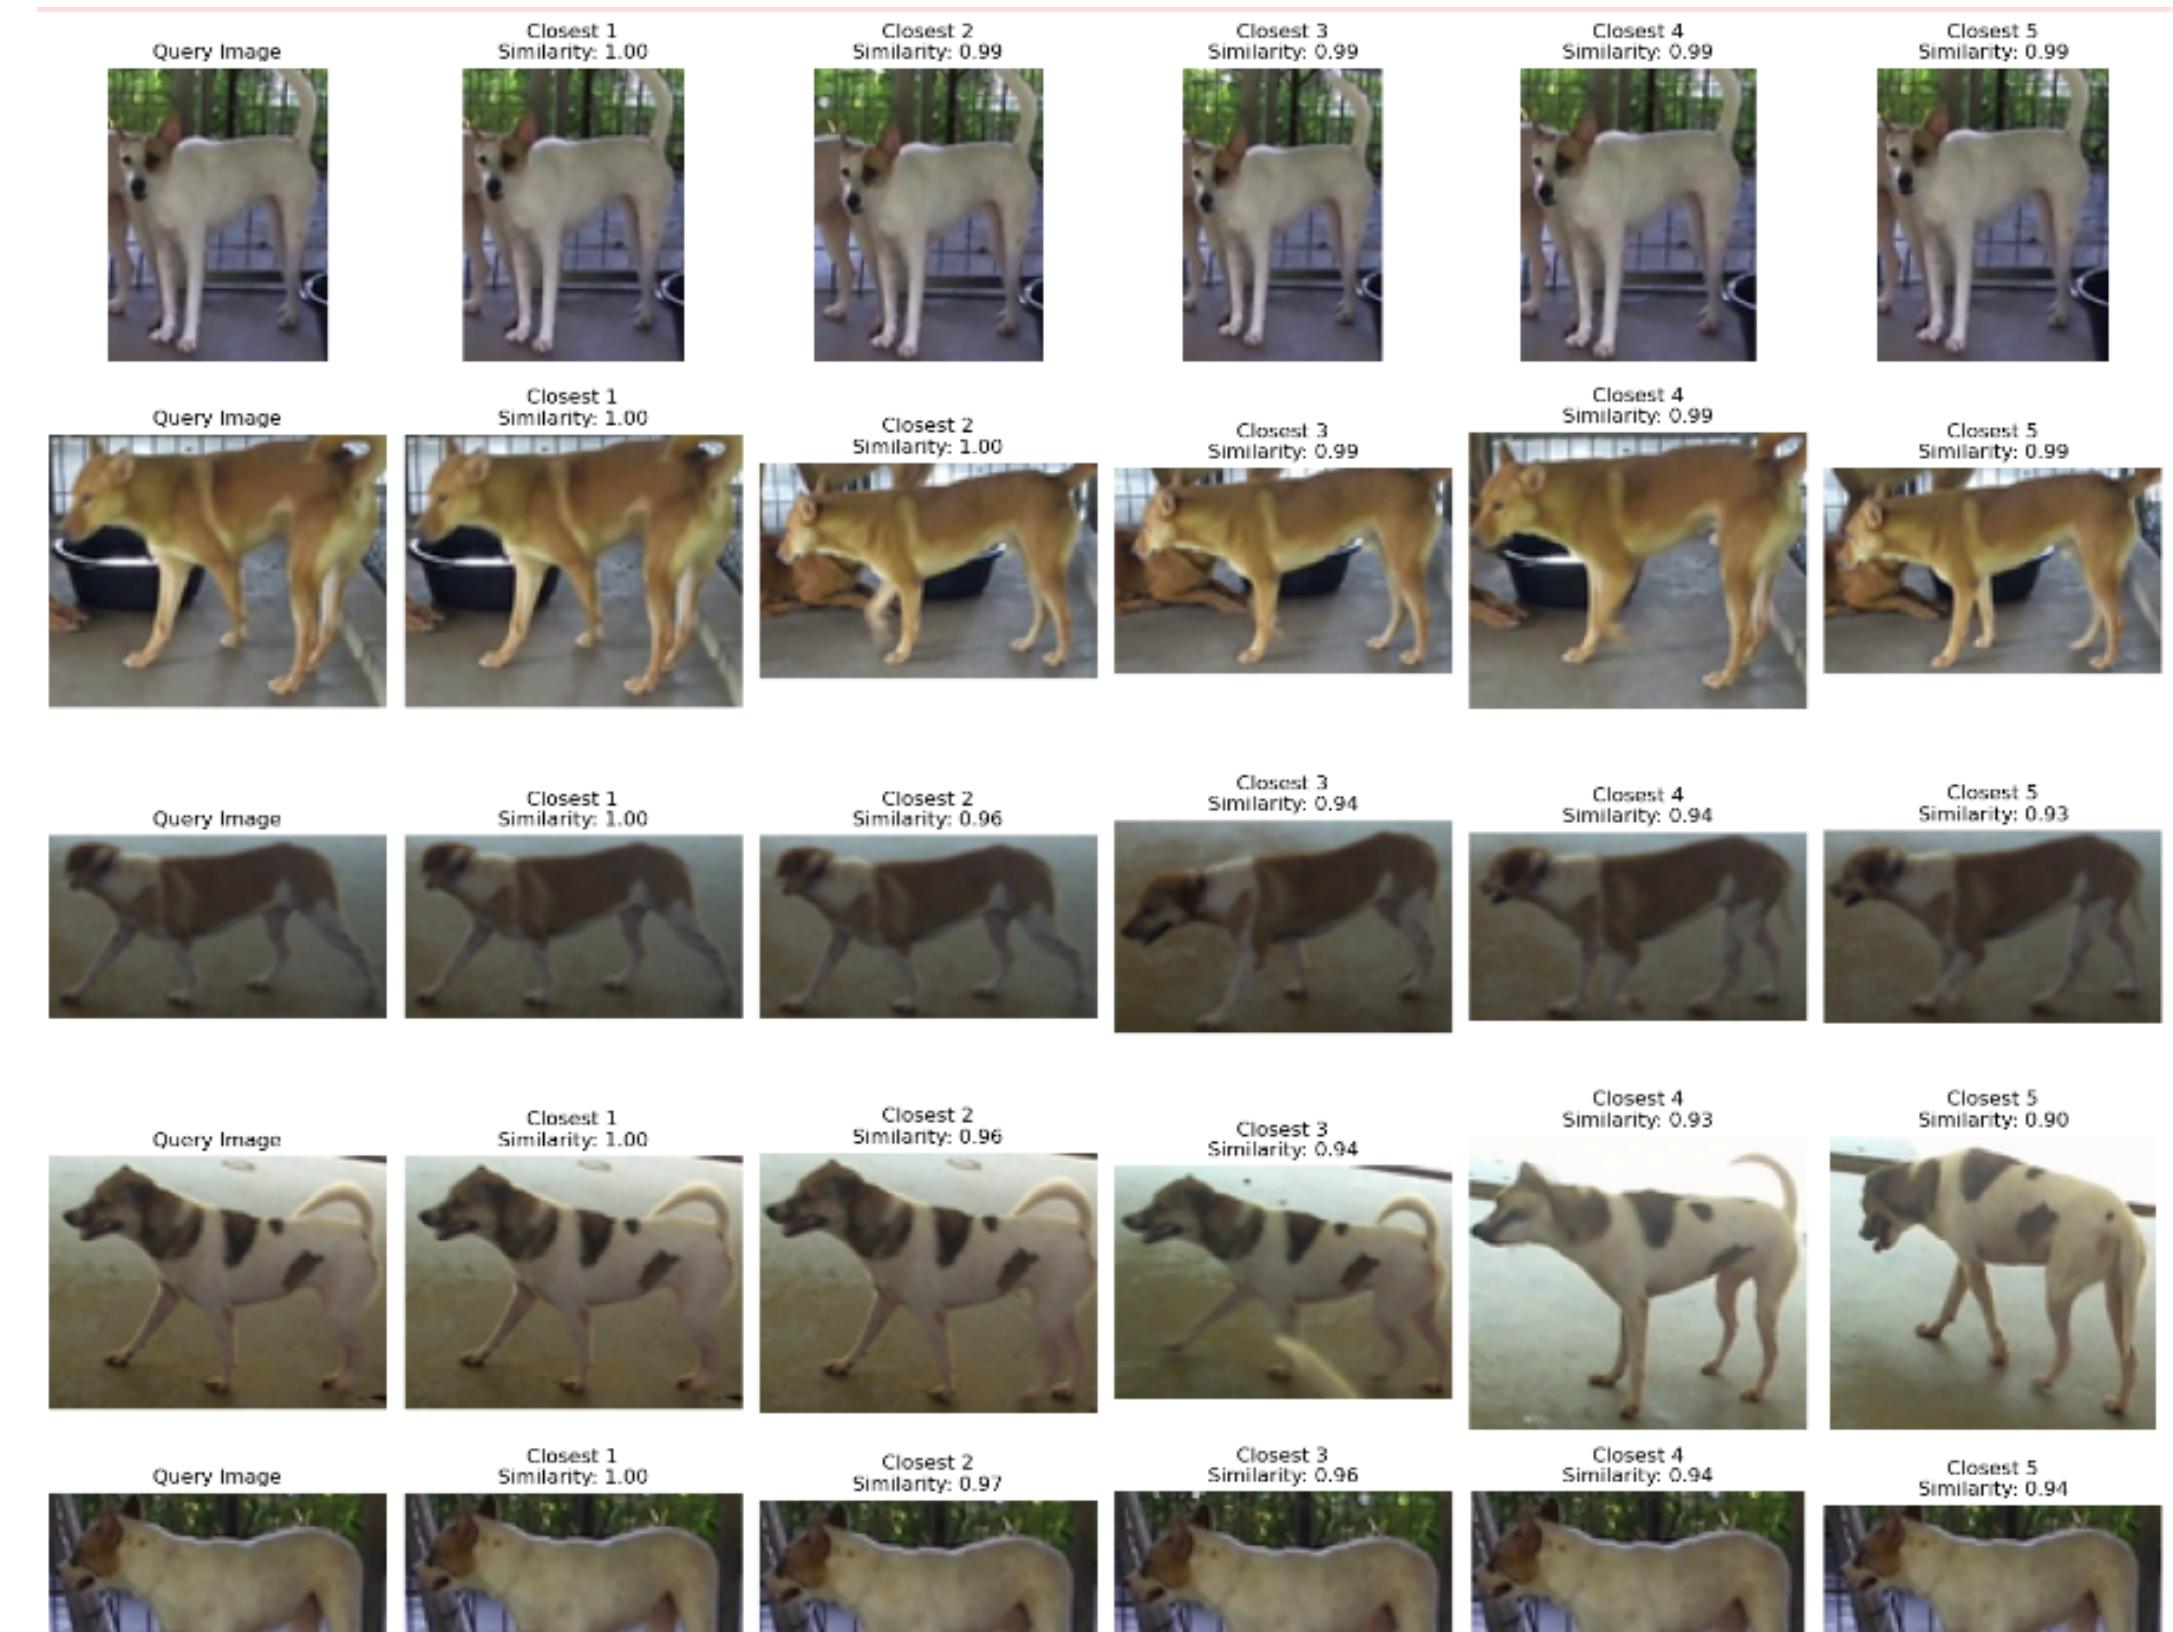
```
print(f'This is the number of trainable weights after freeze some layers {len(embedding_model.trainable_weights)}')
```

This is the number of trainable weights 108

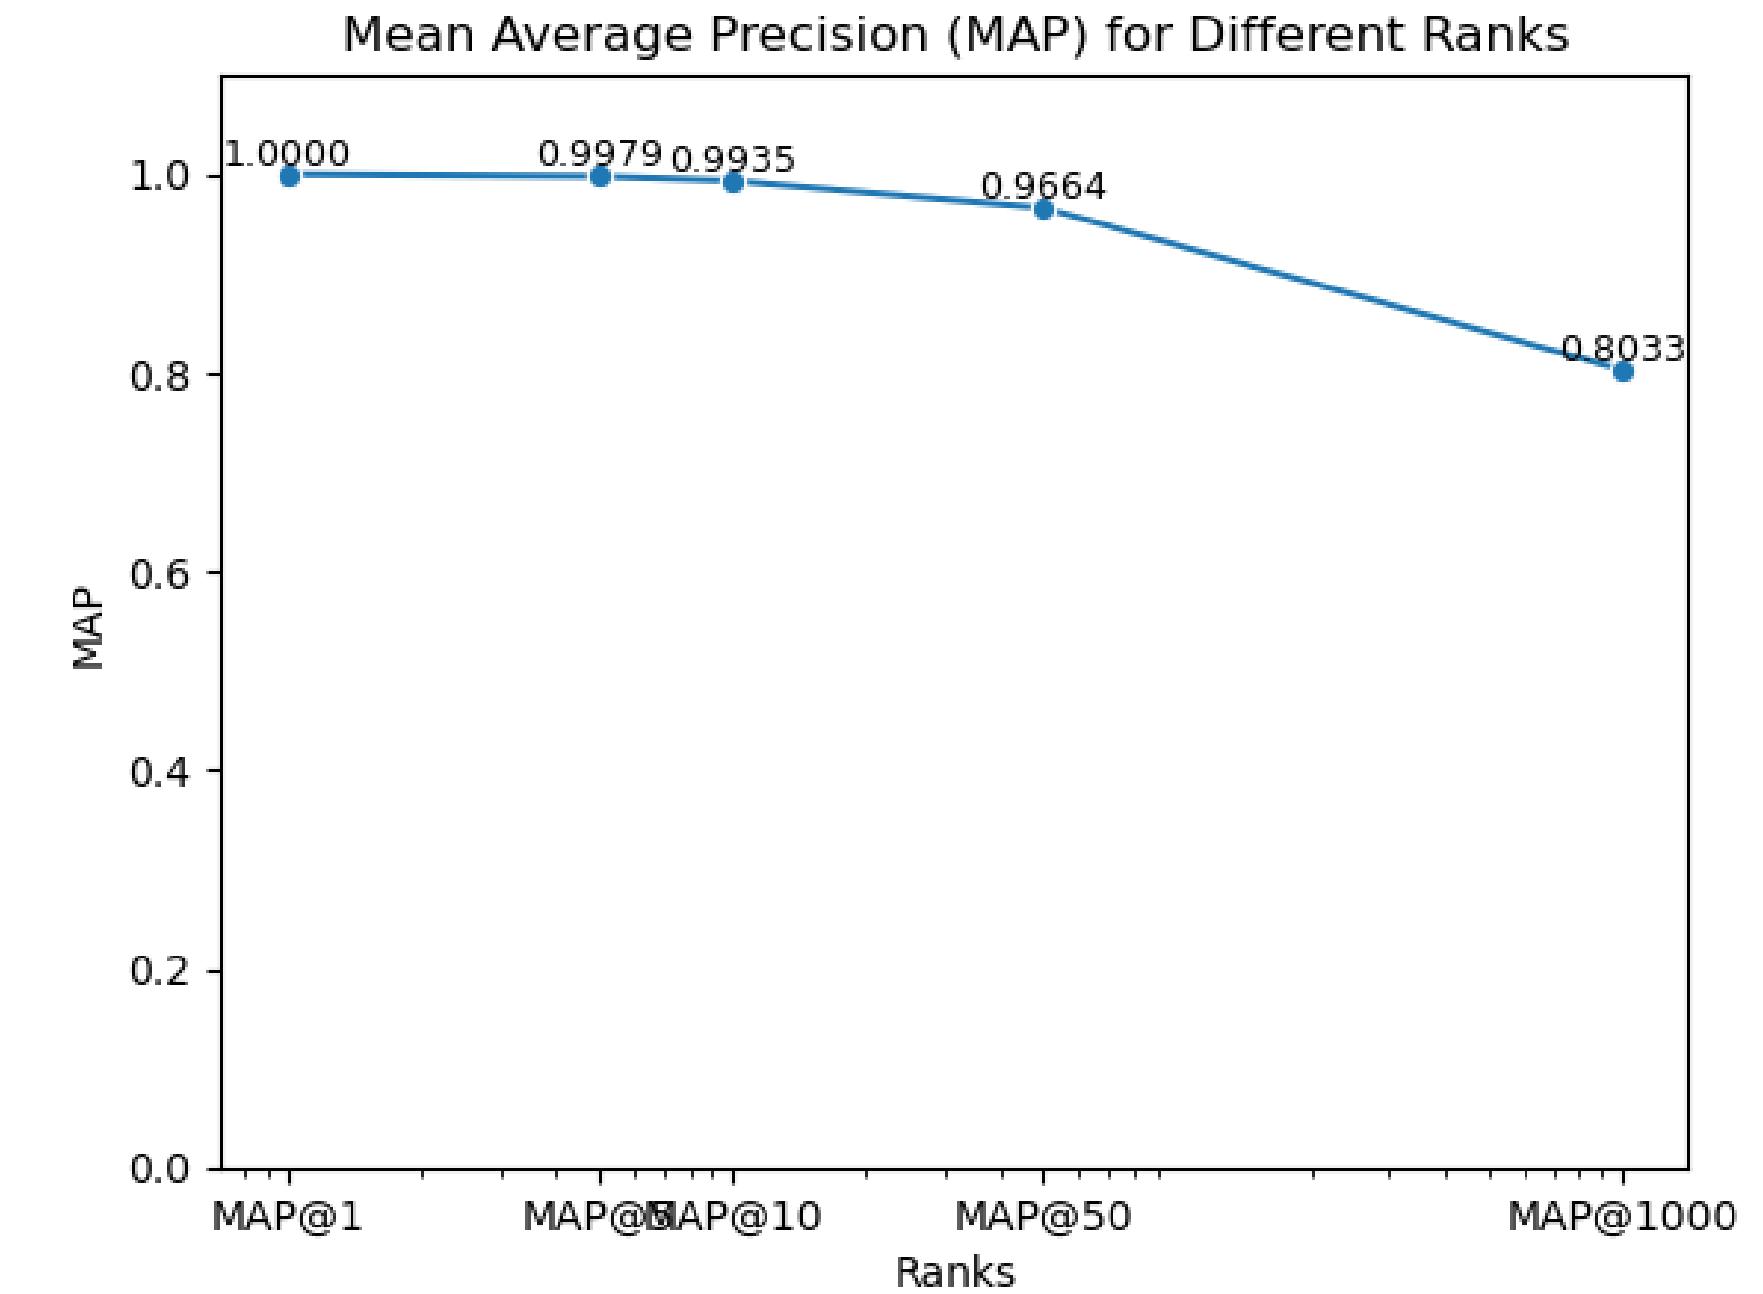
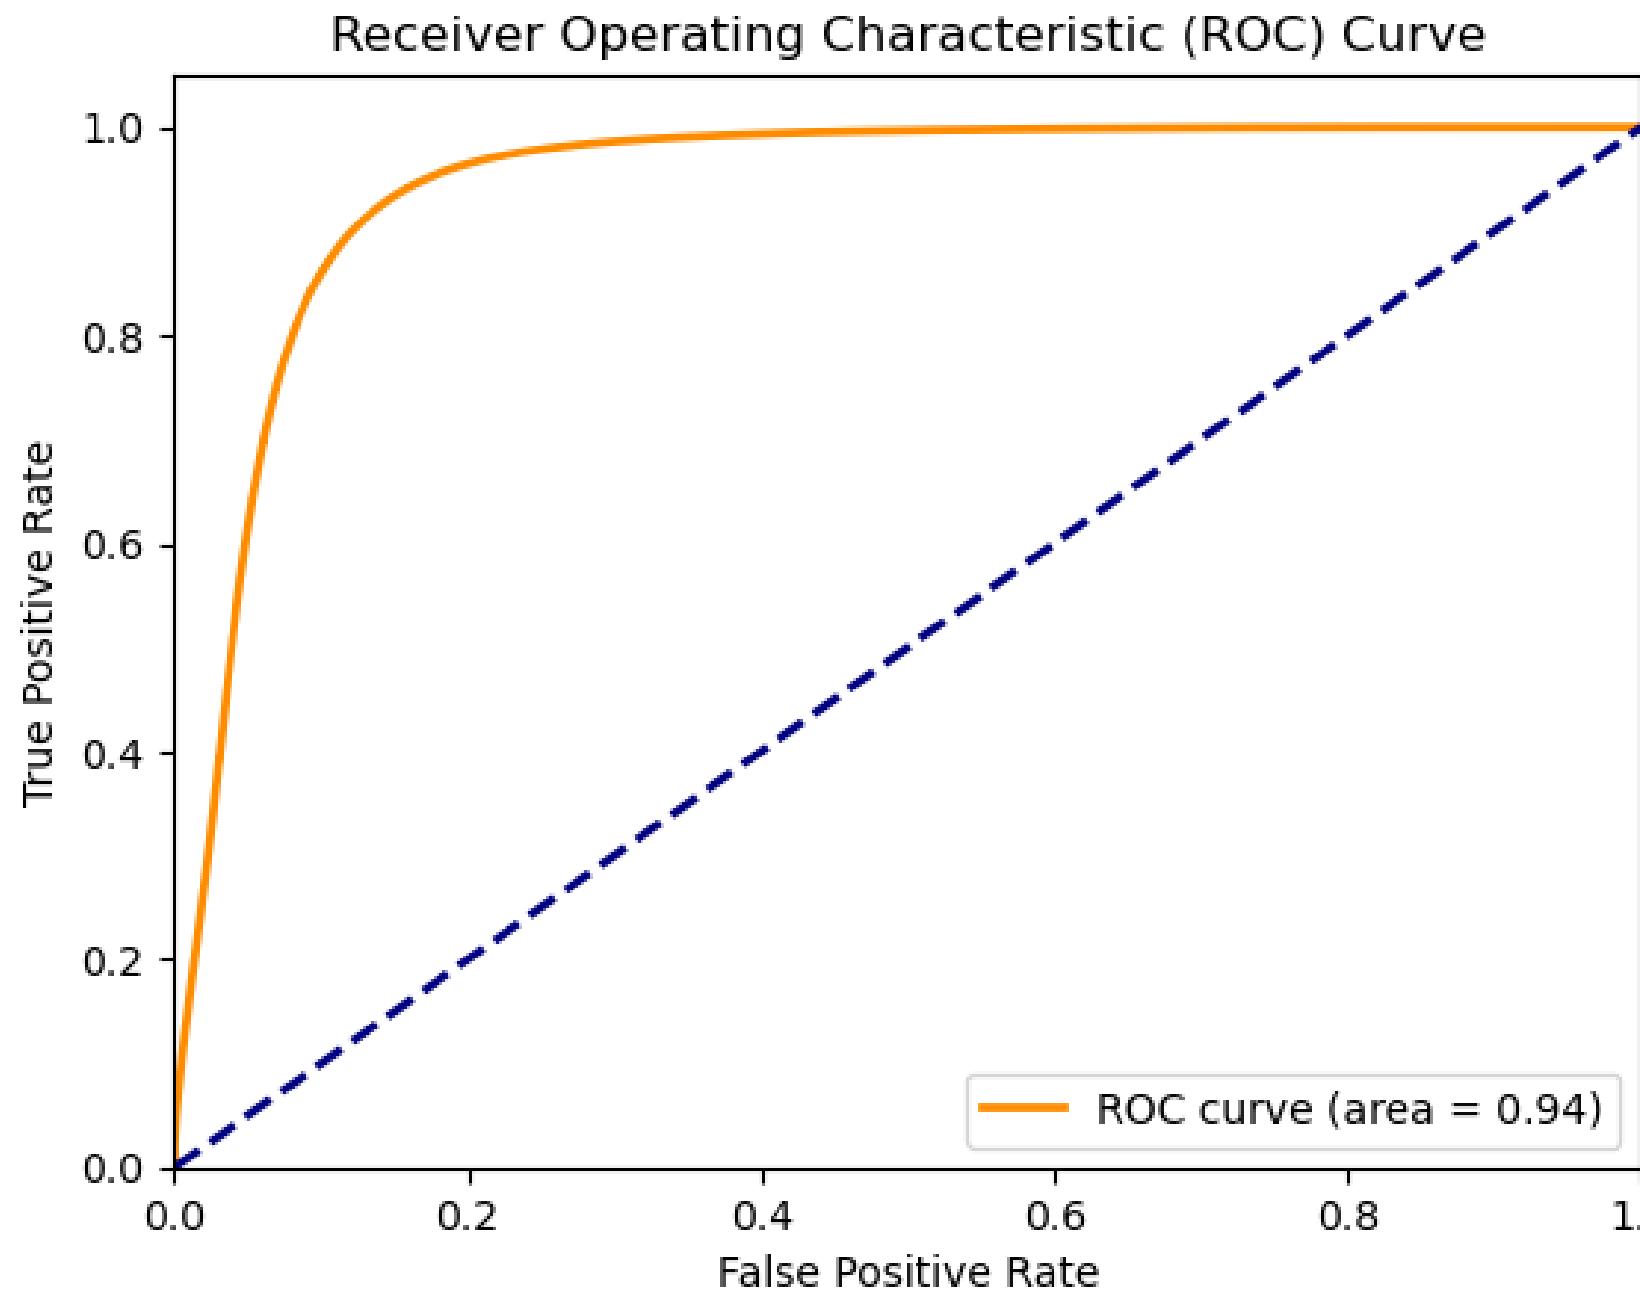
My Siamese Network (inceptionResnetv2)



My Siamese Network (inceptionResnetv2)



My Siamese Network (inceptionResnetv2)



Hyperparameter

```
IMG_SIZE = (224, 224)
BATCH_SIZE = 10
```

```
# Define the triplet loss function
def triplet_loss(_, y_pred):
    margin = 1.0
    anchor_positive_distance, anchor_negative_distance = y_pred[:, 0], y_pred[:, 1]
    loss = tf.maximum(anchor_positive_distance - anchor_negative_distance + margin, 0.0)
    return tf.reduce_mean(loss)

print(f'This is the number of trainable weights {len(embedding_model.trainable_weights)}')
```

This is the number of trainable weights 502

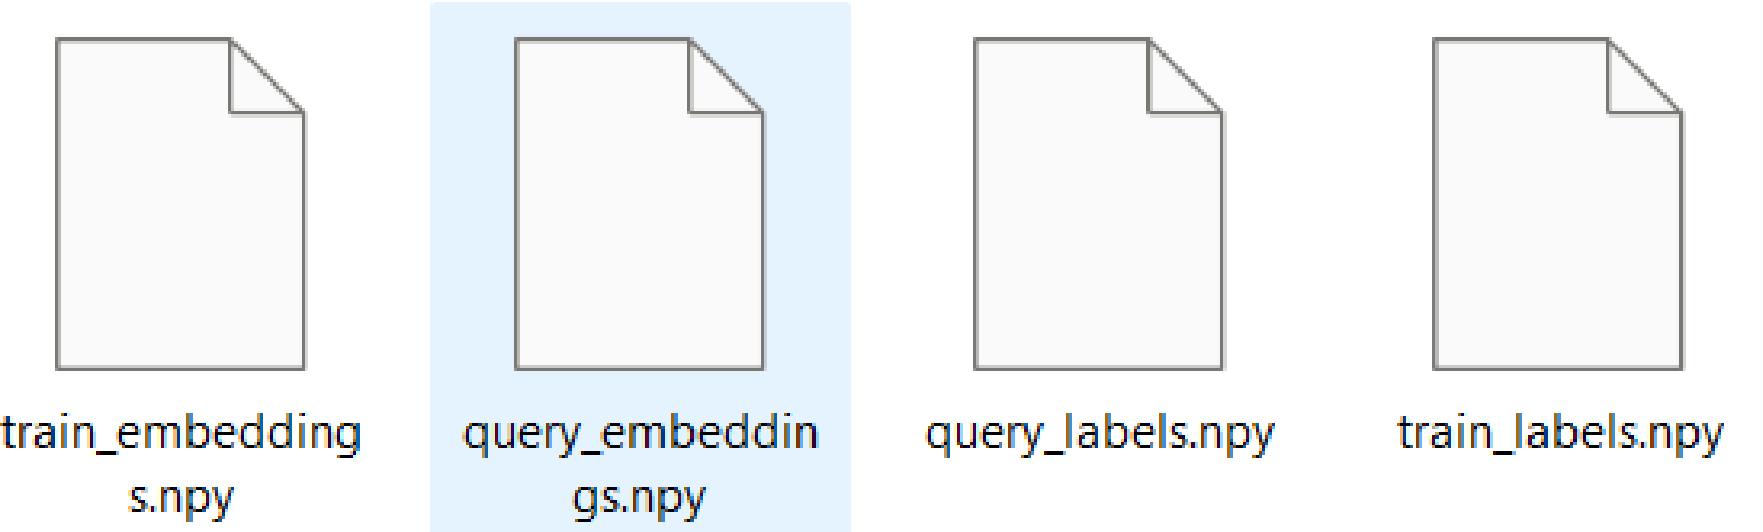
```
from tensorflow.keras.callbacks import EarlyStopping

# Define your early stopping callback
early_stopping = EarlyStopping(monitor='val_loss', patience=5, restore_best_weights=True)

# Train the Siamese model with the learning rate scheduler and early stopping
history = siamese_model.fit(
    train_gen,
    steps_per_epoch=len(train_df) // BATCH_SIZE,
    validation_data=val_gen,
    validation_steps=len(val_df) // BATCH_SIZE,
    epochs=50,
    callbacks=[early_stopping]
)
```

VGG19
VGG16
VGG19
ResNet50
InceptionV3
InceptionResNetV2
MobileNet
MobileNetV2
DenseNet121
DenseNet169
EfficientNetB0

learning rate scheduler with
new data gen



Future work

Hard and Semi-Hard Data Generator:

Instead of storing the embedding features in RAM, save them as .npy files and use garbage collection (gc) after utilizing the data. This approach will be presented in our upcoming presentation.

Ongoing - realtime
use train embedded and query embeded
and train label and query label(not finished)



PongsathornUtsa/siamese_network: Image processing project at KU

Image processing project at KU. Contribute to PongsathornUtsa/siamese_network development by creating an account on GitHub.

 GitHub